

Introduction to 64 Bit Intel Assembly Language Programming

Goals for Cos 284's assembly component

- Learn internal data formats
- Learn basic 64 bit Intel/AMD instructions
- Write pure assembly programs
- Write mixed C/C++ and assembly programs
- Use the gdb debugger for ASM
- Floating point instructions
- Arrays
- Functions
- Structs
- Using system calls, both directly and indirectly
- Data structures and high performance ASM

Problems with assembly language

- Assembly is the poster child for non-portability
 - ▶ Different CPU = different assembly instruction set.
 - ▶ Different OS = different function ABI (application binary interface)
 - ▶ Intel/AMD CPUs operate in 16, 32, and 64 bit modes
- Difficult to program
 - ▶ More time = more money
 - ▶ Less reliable as it is very error prone.
 - ▶ Difficult to maintain
- Syntax does not resemble mathematics
- No syntactic protection
 - ▶ No structured ifs, loops
- No typed variables
 - ▶ Can use a pointer as a floating point number
 - ▶ Can load a 4 byte integer from a double variable
- Variable access is roughly like using pointers

What's good about assembly language?

- Assembly language is fast
 - ▶ Optimizing C/C++ compilers will be faster than a novice most of the time.
 - ▶ You need to dissect an algorithm and rearrange it to use a special feature that the compiler can't figure out
 - ▶ Generally you must use a special instructions
 - ▶ There are over 1000 instructions
 - ▶ Still it can be faster
- Assembly programs are small
 - ▶ But memory is cheap and plentiful
 - ▶ C/C++ compilers can optimize for size
 - ▶ Compilers can re-order code sections to reduce size
- Assembly can do things not possible in C/C++
 - ▶ I/O instructions
 - ▶ Manage memory mapping registers
 - ▶ Manipulate other internal control registers

What's good about assembly for ordinary mortals?

- Teaches you how the programs really works
- Teaches you how storage and arithmetic is done in registers
- Teaches you C/C++ function register and stack usages
- Teaches you how stack frames are built and destroyed.
- Optimization techniques are explained.
- Computer bugs are more immediately related to machine instructions and limitations
- You will learn how the compiler implements
 - ▶ if/else statements
 - ▶ loops
 - ▶ functions
 - ▶ structures
 - ▶ arrays
 - ▶ recursion
- Your coding will improve.

Generation of languages

- First generation - machine language
- Second generation - assembly language
 - ▶ Names for instructions
 - ▶ Names for variables
 - ▶ Names for locations of instructions
 - ▶ Perhaps with macros - code replacement
- Third generation - not machine instructions
 - ▶ Modeled after mathematics - Fortran
 - ▶ Modeled after English - Cobol
 - ▶ List processing - Lisp
- Fourth generation - domain specific
 - ▶ SQL
- Fifth generation - describe problem, computer generates algorithm
 - ▶ Prolog

Assembly example

```
; Program: exit
;
; Executes the exit system call
;
; No input
;
; Output: only the exit status ($? in the shell)
;
segment .text
global _start
_start:
    mov     eax,1          ; 1 is the exit syscall number
    mov     ebx,5          ; the status value to return
    int     0x80           ; execute a system call
```

Assembly syntax

- ; starts comments
- Labels are strings which are not instructions
 - ▶ Usually start in column 1
 - ▶ Can end with a colon to avoid confusion with instructions
- Instructions can be machine instructions or assembler instructions
 - ▶ `mov` and `int` are machine instructions or opcodes
 - ▶ `segment` and `global` are assembler instructions or pseudo-ops
- Instructions can have operands
 - ▶ `here: mov eax, 1`
 - ▶ `here` is a label for the instruction
 - ▶ `mov` is an opcode
 - ▶ `eax` and `1` are operands

Some assembler instructions

- `section` or `segment` define a part of the program
 - ▶ `.text` is where instructions go for Linux
- `global` defines a label to be used by the linker
- `global _start` makes `_start` a global label
- `_start` or `main` is where a program starts
 - ▶ `_start` is more basic
 - ▶ `main` is called (perhaps indirectly) by `_start`

Assembling the exit program

- `yasm -f elf64 -g dwarf2 -l exit.lst exit.asm`
- `-f elf64` says we want a 64 bit object file (elf=extensible linking format)
- `-g dwarf2` says we want dwarf2 debugging info (why dwarf?)
 - ▶ dwarf2 works pretty well with the gdb debugger
- `-l exit.lst` asks for a listing in `exit.lst`
- `yasm` will produce `exit.o`, an object file
 - ▶ machine instructions not ready to execute

exit.lst

```
1                                     %line 1+1 exit.asm
2
3
4
5
6
7
8
9
10                                  [segment .text]
11                                  [global _start]
12
13                                  _start:
14 00000000 B801000000              mov eax,1
15 00000005 BB05000000              mov ebx,5
16 0000000A CD80                   int 0x80
```

Linking

- Linking means combining object files to make an executable file
- For programs with `_start`
 - ▶ `ld -o exit exit.o`
 - ▶ Builds a file named `exit`
 - ▶ Default is `a.out`
- For programs with `main`
 - ▶ `gcc -o exit exit.o`
 - ▶ Gets default `_start` function from the C library
- `./exit` to run the program

Floating point numbers

Consider 1.75, in 32bit-IEEE 754 the number becomes:

0	01111111	110000000000000000000000
Positive	127	(1).75

Grouping into 4 bit nibbles:

0011	1111	1110	0000	0000	0000	0000	0000
3	f	e	0	0	0	0	0

But this is stored reversed and with each nibble pair swapped:

0	0	0	0	e	0	3	f
---	---	---	---	---	---	---	---

listings example part 1

Consider the following asm file “fp.asm”.

```
1      segment .data
2 zero    dd      0.0
3 one     dd      1.0
4 neg1    dd      -1.0
5 a       dd      1.75
6 b       dd      122.5
7 d       dd      1.1
8 e       dd      10000000000.0
```

The **dd** command specifies a double word data item. A word is 2 bytes.
So a double word is 32 bits.

- **dw** is a data word
- **db** is a byte
- **dq** is a data quad-word

listings example part 1

Now if we create the file listing using:

```
yasm -f elf64 -g dwarf2 -l fp.lst fp.asm
```

The result is:

1	%line 1+1 fp.asm
2	[section .data]
3 00000000 00000000	zero dd 0.0
4 00000004 0000803F	one dd 1.0
5 00000008 000080BF	neg1 dd -1.0
6 0000000C 0000E03F	a dd 1.75
7 00000010 0000F542	b dd 122.5
8 00000014 CDCC8C3F	d dd 1.1
9 00000018 F9021550	e dd 10000000000.0

Memory mapping

- Computer memory is an array of bytes from 0 to $n - 1$ where n is the memory size
- Programs perceive “logical” addresses which are mapped to physical addresses
- 2 people can run a program starting at logical address 0x4004c8 while using different physical memory
- CPU translates logical addresses to physical during instruction execution
- The CPU translation can be just as fast as if the software used physical addresses
- The x86-64 CPUs can map pages of sizes 4096 bytes and 2 megabytes
- Linux uses 2 MB pages for the kernel and 4 KB pages for programs
- Some recent CPUs support 1 GB pages

Translating an address

- Suppose an instruction references address 0x43215628
- With 4 KB pages, the rightmost 12 bits are an offset into a page
- With 0x43215628 the page offset is 0x628
- The page number is 0x43215
- Let's assume that the computer is set up to translate page 0x43215 to physical addresses 0x7893000 - 0x7893fff
- Then address 0x43215628 is mapped to 0x7893628

Benefits of memory mapping

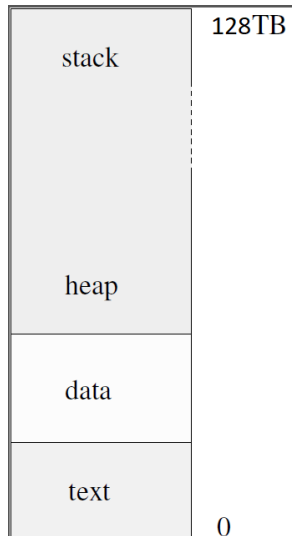
- User processes are protected from each other
 - ▶ Your process can't read my process's data
 - ▶ Your process can't write my data
- The operating system is protected from malicious or errant code
- It is easy for the operating system to give processes contiguous chunks of “logical” memory

Why study memory mapping?

- If you write programs, the mapping is automatic
- We will not discuss instructions for changing mapping tables
- So what difference does it make?
- It helps explain page faults
 - ▶ Suppose you allocate an array of 256 bytes at logical address 0x45678200
 - ▶ Then all addresses from 0x45678000 to 0x45678fff are valid
 - ▶ You can go well past the end of the array before you can get a segmentation violation
- Knowledge is power!

Process memory model in Linux

- A Linux process has 4 logical segments
 - ▶ text: machine instructions
 - ▶ data: static data initialized when the program starts
 - ▶ heap: data allocated by `malloc` or `new`
 - ▶ stack: run-time stack
 - ★ return addresses
 - ★ some function parameters
 - ★ local variables for functions
 - ★ space for temporaries
- In reality it is slightly more complex
- 128TB is 47 bits of all 1's (≈ 141 base 10, TB)
- CPU could use 48 bit logical addresses



Memory segments

- The text segment is named `.text` in `yasm`
 - ▶ `_start` and `main` are not actually at 0
 - ▶ The text segment does not need to grow, so the data segment can be placed immediately after it
- The data segment is in 2 parts
 - ▶ `.data` which contains initialized data
 - ▶ `.bss` which contains reserved data
 - ▶ “bss” stands for “Block Started by Symbol”
- The heap and the stack both need to grow
 - ▶ The heap grows up
 - ▶ The stack grows down
 - ▶ They meet in the middle and explode

Stack segment limits

- The stack segment is limited by the Linux kernel
- The typical size is 16 MB for 64 bit Linux
- This can be inspected using “`ulimit -a`” or “`ulimit -s`”
- 16 MB seems fairly small, but it is fine until you start using large arrays as local variables in functions
- The stack address range is 0x7fffffff000000 to 0xffffffffffff
- A fault to addresses in this range are recognized by the kernel to allow the stack to grow as needed

Memory example source code

```

        segment .data
a        dd      4
b        dd      4.4
c        times   10 dd 0
d        dw      1, 2
e        db      0xfb
f        db      "hello world", 0

        segment .bss
g        resd     1
h        resd     10
i        resb     100
```

Memory example listing file

1		%line 1+1 memory.asm
2		[section .data]
3	00000000 04000000	a dd 4
4	00000004 CDCC8C40	b dd 4.4
5	00000008 00000000<rept>	c times 10 dd 0
6	00000030 01000200	d dw 1, 2
7	00000034 FB	e db 0xfb
8	00000035 68656C6C6F20776F72-	f db "hello world", 0
9	00000035 6C6400	

- Addresses are relative to start of .data in this file
- $b = 0x408ccccd = 0\ 10000001\ 00011001100110011001101$
- Sign bit is 0, exponent field is $0x81 = 129$, \implies actual exponent = 2
- Fraction is $1.00011001100110011001101$

Memory example listing file (2)

```
11                                [section .bss]
12 00000000 <gap>                g resd 1
13 00000004 <gap>                h resd 10
14 0000002C <gap>                i resb 100
```

- Notice that the addresses start again at 0
- The commands reserve space
- `resd 1` reserves 1 double word or 4 bytes
- `resd 10` reserves 10 double words or 40 bytes
- `resb 100` reserves 100 bytes

Examining memory

Useful tools to examine memory are:

- gdb
- ebe

Memory Mapping in 64 Bit Mode and Registers

Chapter 4


- What is page?

- What is page?
 - ▶ A page is a fixed-length contiguous block of virtual memory, described by a single entry in the page table.
 - ▶ It is the smallest (usually) unit of data for memory management in a virtual memory operating system.
 - ▶ In our case we will consider 4kb pages.

- What is page?
 - ▶ A page is a fixed-length contiguous block of virtual memory, described by a single entry in the page table.
 - ▶ It is the smallest (usually) unit of data for memory management in a virtual memory operating system.
 - ▶ In our case we will consider 4kb pages.
 - ▶ Check with `getconf PAGESIZE`.
- How is the mapping of a page between virtual address space and physical address space done?

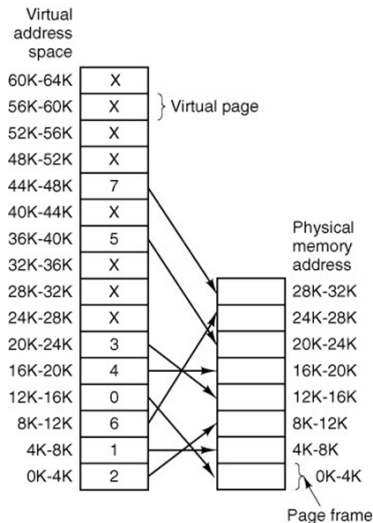
Single level paging

- To the right we have an example of a single layer paging system
- The X's indicate that the virtual page does not at present have a corresponding physical page.
 - ▶ If a program requires a virtual page marked with an X this causes a page fault.
 - ▶ The system must then allocate some physical memory to assign this virtual page to.
 - ★ But what if there is no space left?




Single level paging

- The primary problems with a single level paging system is the size of the table.
 - ▶ Consider the linux process model with a virtual memory space of 128TB that would imply a table with $2^{(47-12)} = 2^{35}$ entries.
 - ▶ About 34 billion entries!
 - ▶ Assuming a 1GB of physical memory we would need to address $2^{(30-12)} = 2^{18}$ physical 4k pages.
 - ★ So we would need 18-bits plus 1 bit to indicate if the entry is valid or invalid.
 - ▶ Just this table will occupy $2^{35} * 19$ bits, which is about 81.6 GB.




Multilevel Page Tables

- To get around the problem of having to store huge page tables in memory all the time, many computers use a multilevel page table.
- As a simple example consider we have a 32-bit virtual address (4GB), that is partitioned into a
 - ▶ 10-bit PT1 field
 - ▶ 10-bit PT2 field,
 - ▶ and a 12-bit Offset field (for the 4k pages).
- In general if we used PT1+PT2 together we would be working 2^{20} pages
 - ▶ That's a lot of pages to keep in memory




Multilevel Page Tables

- The **top** level page table corresponds to the first 10-bits.
 - ▶ which maps to 1024 page tables
- The **second** level page table corresponds to the second 10-bits
 - ▶ Which maps to the 1024 pages of size 4K
- The last remaining 12 bits are used as the offset to address the contents of the 4K page.



Multilevel Page Tables

- The secret to the multilevel page table method is to avoid keeping all the page tables in memory all the time.
 - ▶ In particular, those that are not needed should not be kept around.
- By marking elements of the top-level page table as absent we do not need to maintain (or store) all the second level page tables, saving substantially on space.



Memory mapping pages and tables

- Linux uses 4 layer page tabling.
- Each page is $2^{12} = 4096$ bytes
- An address is 8 bytes (not all used)
- Each page table can hold $2^9 = 512$ addresses
- A 9 bit field is needed to index the mapping tables
- Current mapping uses 48 bits, so we are limited to 2^{48} bytes which is 256 TB

Logical memory address fields

63-48	47-39	38-30	29-21	20-12	11-0
unused	PML4 index	page directory pointer index	page directory index	page table index	page offset

- Bits 47-39 are used to index the PML4 table
- Bits 38-30 are used to index the selected page directory pointer table
- Bits 29-21 are used to index the selected page directory table
- Bits 20-12 are used to index the selected page table
- Bits 11-0 are the offset into the page (for 4 KB pages)

Large pages

- Using the first 3 existing levels of page tables, we can have large pages with $2^{21} = 2097152$ bytes.
- This is used by Linux for the kernel

CPU support for fast lookups

- A CPU uses a special cache called a “Translation Lookaside Buffer” or TLB to speed up memory translation
- A TLB operates much like a hash table
- Presented with a logical address, it produces the physical address or failure in about 1/2 a clock cycle
- The Intel Core i7 has 2 levels of TLBs
 - ▶ Level 1 holds 64 small page translations (or 32 big pages)
 - ▶ Level 2 holds 512 page translations
 - ▶ Large programs with small pages will experience TLB misses which can be satisfied fairly rapidly with normal cache
 - ▶ Very large programs can crawl

Chapter 5

Register basics

- Computer main memory has a latency of about 80 nanoseconds
- A 3.3 GHz CPU uses approximately 0.3 nsecs per cycle
- Memory latency is about 240 cycles
- The Core i7 has 3 levels of cache with different latencies
 - ▶ Level 3 about 48 nsec latency or about 150 cycles
 - ▶ Level 2 about 10 nsec latency or about 39 cycles
 - ▶ Level 1 about 4 nsec latency or about 12 cycles
- There is a need for even faster memory
- This ultra-fast “memory” is the CPU’s registers
- Some register-register instructions complete in 1 cycle

x86-64 registers

- CPUs running in x86-64 mode have 16 general purpose registers
- There are also 16 floating point registers (XMM0-XMM15)
- There is also a floating point register stack which we ignore
- The general purpose registers hold 64 bits
- The floating point registers can be either 128 or 256 bits
 - ▶ The CPU can use them to do 1 32 bit or 1 64 bit floating point operation in an instruction
 - ▶ The CPU can also use these to do packed operations on multiple integer or floating point values in an instruction
 - ▶ “Single Instruction Multiple Data” - SIMD
- The CPU has a 64 bit instruction pointer register - `rip`
 - ▶ contains the address of the next instruction to execute.
- There is a 64 bit flags register, `rflags`, holding status values like whether the last comparison was positive, zero or negative

General purpose registers

- These registers evolved from 16 bit CPUs to 32 bit mode and finally 64 bit mode
- Each advance has maintained compatibility with the old instructions
- The old register names still work
- The old collection was 8 registers which were not entirely general purpose
- The 64 bit collection added 8 completely general purpose 64 bit registers named r8 - r15

The 64 bit registers evolved from the original 8

- Software uses the “r” names for 64 bit use, the “e” names for 32 bit use and the original names for 16 bit use
- rax - general purpose, accumulator
 - ▶ rax uses all 64 bits
 - ▶ eax uses the low 32 bits
 - ▶ ax uses the low 16 bits
- rbx, ebx, bx - general purpose
- rcx, ecx, cx - general purpose, count register
- rdx, edx, dx - general purpose
- rdi, edi, di - general purpose, destination index
- rsi, esi, si - general purpose, source index
- rbp, ebp, bp - general purpose, stack frame base pointer
- rsp, esp, sp - stack pointer, rsp is used to push and pop


The original 8 registers as bytes

- Kept from the 16 bit mode
 - ▶ `al` is the low byte of `ax`, `ah` is the high byte
 - ▶ `bx` can be used as `bl` and `bh`
 - ▶ `cx` can be used as `cl` and `ch`
 - ▶ `dx` can be used as `dl` and `dh`
- New to x86-64
 - ▶ `dil` for low byte of `rdi`
 - ▶ `sil` for low byte of `rsi`
 - ▶ `bpl` for low byte of `rbp`
 - ▶ `spl` for low byte of `rsp`
- There is no special direct way to access any “higher” bytes of registers

Visual summary

Break down of a 64-bit register.

Summary



The 8 new general purpose registers as smaller registers

- Here the naming convention changes
- Appending “d” to a register accesses its low double word - r8d
 - ▶ double word = 4 bytes = 32 bits.
- Appending “w” to a register accesses its low word - r12w
 - ▶ single word = 2 bytes = 16 bits.
- Appending “b” to a register accesses its low byte - r15b

Moving a constant into a register

- Moving is fundamental
- `yasm` uses the mnemonic `mov` for all sorts of moves
- The generated code from `gcc` uses mnemonics like `movq`
 - ▶ `gcc` uses AT&T syntax by default
- Most instructions can use 1, 2 or 4 byte immediate fields
- `mov` can use an 8 byte immediate value.

```
mov rax, 0x0123456789abcdef ; can move 8 byte immediates
mov rax, 0
mov eax, 5                   ; the upper half is set to 0
mov r8w, 16                  ; affects only low word
```

Moving a value from memory into a register

```
segment .data
a      dq      175
b      dq      4097
c      db      1, 2, 3, 4, 5, 6, 7, 8
d      dd      0xffffffff

segment .text
mov     rax, a
mov     rbx, [a]
mov     rcx, [c]
mov     edx, [d]
```

- Using simply **a** places the address of **a** into **rax**
- Using **[a]** places the value of **a** into **rbx**

A program to add 2 numbers from memory

```
        segment .data
a        dq      175
b        dq      4097
        segment .text
        global  main
main:
        mov     rax, [a]      ; mov a into rax
        add     rax, [b]      ; add b to rax
        xor     rax, rax
        ret
```

Move with sign extend or zero extend

- If you move a double word into a double word register (e.g. `eax`), the upper half is zeroed out
- If you move a 32 bit immediate into a 64 bit register it is sign extended
- Sometimes you might wish to load a smaller value from memory and fill the rest of the register with zeroes
- Or you may wish to sign extend a small value from memory
- For `movsx` and `movzx` you need a size qualifier for the memory operand

```
movsx  rax, byte [data]      ; move byte, sign extend
movzx  rbx, word [sum]       ; move word, zero extend
movsxd rcx, dword [count]    ; move dword, sign extend
```

Moving values from a register into memory

- Simply use the memory reference as the first operand

```
mov    [a], rax        ; move a quad word to a
mov    [b], ebx        ; move a double word to b
mov    [c], r8w        ; move a word to c
mov    [d], r15b       ; move a byte to d
```

Moving data from one register to another

- Use 2 register operands

```
mov    rax, rbx    ; move rbx to rax
mov    eax, ecx    ; move ecx to eax, zero filled
mov    cl, al      ; move al to cl, leave rest of
; unchanged
```

A Little Bit of Math

A little bit of math

- So far we have learned how to get values into registers
- And how to place them back into memory
- Just some ordinary arithmetic can help us write slightly more useful programs
- We will only discuss integer math in this lecture.

Negation

- The negate instruction, `neg`, converts a number to its **two's complement**.
- `neg` sets the sign and zero flags
 - ▶ Which will be useful when we perform conditional jumps and moves.
- There is only a single operand which is source and destination

```
neg    rax    ; negate the value in rax
neg    eax    ; negate the value of eax and zx the rest
neg    ax     ; negate the value of ax
neg    al     ; negate the value of al
```

Negation

- For memory operands you must include a size prefix
- The sizes are byte, word, dword and qword

```
neg    qword [x] ; negate a 8 byte integer at x
neg    dword [x] ; negate a 4 byte integer at x
neg    word  [x] ; negate a 2 byte integer at x
neg    byte  [x] ; negate a 1 byte integer at x
```

The add instruction

- The add instruction always has exactly 2 operands
 - ▶ The source and, (RHS)
 - ▶ the destination (LHS)
- It adds its source value to its destination
- The source can be a
 - ▶ immediate
 - ▶ register
 - ▶ memory location
- The destination can be a
 - ▶ register
 - ▶ memory location
- Using memory locations for both source and destination is **not allowed**
 - ▶ as is the general pattern with the x86-64 instruction set.

The add instruction

- After an ADD instruction executes it sets the following flags:
 - ▶ sign flag(SF)
 - ▶ zero flag(ZF)
 - ▶ overflow flag(OF)
 - ▶ there are more, but they are no important in this course.
- There is no special “signed add” versus “unsigned add” since the logic is identical
- There is a special 1 operand increment instruction, `inc`

```
inc rax      ; add one to rax
```

```
inc byte [x]; add one to the interger byte at x
```

A program using add

Program has three “variables”: **a**=151, **b**=310, and **sum**=0. We want to:

- set **a**=**a**+9
- set **sum**=**a**+**b**+10

A program using add

```
segment .data
a      dq      151
b      dq      310
sum     dq      0
segment .text
global  main

main:
    mov     rax, 9          ; set rax to 9
    add     [a], rax        ; add rax to a
    mov     rax, [b]        ; get b into rax
    add     rax, 10         ; add 10 to rax
    add     rax, [a]        ; add the contents of a
    mov     [sum], rax      ; save the sum in sum
    xor     rax, rax
    ret
```

The subtract instruction

- The sub instruction performs integer subtraction
- Like add it supports 2 operands
- Only one of the operands can be a memory operand
- There is a “subtract one” instruction, dec
- It sets the sign flag, the zero flag and the overflow flag
- There is no special “signed subtract” versus “unsigned subtract” since the logic is identical

A program using sub

Program has three “variables”: **a**=100, **b**=200, and **diff**=0. We want to:

- set **a**=**a**-10
- set **b**=**b**-10
- set **diff**=**b**-**a**

A program using sub

```
segment .data
a      dq      100
b      dq      200
diff   dq      0
segment .text
global main

main:
    mov     rax, 10
    sub     [a], rax      ; subtract 10 from a
    sub     [b], rax      ; subtract 10 from b
    mov     rax, [b]      ; move b into rax
    sub     rax, [a]      ; set rax to b-a
    mov     [diff], rax  ; move the difference to diff
    mov     rax, 0
    ret
```

Multiplication

- Unsigned multiplication is done using the `mul` instruction
- Signed multiplication is done using `imul`
- There is only 1 form for `mul`
 - ▶ It uses 1 operand, the source operand
 - ▶ The other factor is in `rax`, `eax`, `ax` or `al`
 - ▶ The destination is `ax` for byte multiplies
 - ▶ Otherwise the product is in `rdx:rax`, `edx:eax`, or `dx:ax`

```
mov    rax, [a]
mul    qword [b]    ; a * b will be in rdx:rax
mov    eax, [c]
mul    dword [d]    ; c * d will be in edx:eax
```

Signed multiplication

- `imul` has a single operand form just like `mul`
- It also has a 2 operand form, source and destination, like `add` and `sub`
- Finally there is a 3 operand form: destination, source and immediate source
- If you need all 128 bits of product, use the single operand form

```
imul    rax, 100           ; multiply rax by 100
imul    r8, [x]            ; multiply r8 by x
imul    r9, r10            ; multiply r9 by r10
imul    r8, r9, 11         ; store r9 * 11 in r8
```

Division

- Division returns a quotient and a remainder
- It also has signed (`idiv`) and unsigned forms (`div`)
- In both forms the dividend is stored in `rdx:rax` or parts thereof
- The quotient is stored in `rax`
- The remainder is stored in `rdx`
- No flags are set

```
mov     rax, [x]           ; x will be the dividend
mov     rdx, 0              ; 0 out rdx, so rdx:rax == rax
idiv    qword [y]          ; divide by y
mov     [quot], rax        ; store the quotient
mov     [rem], rdx         ; store the remainder
```

Conditional move instructions

- There are many variants of conditional move, `cmovCC`, where CC is a condition like **I** to mean **less**
- These are great for simple conditionals
- You can avoid interrupting the instruction pipeline

Instruction	effect
<code>cmovz</code>	move if zero flag set
<code>cmovnz</code>	move if zero flag not set (not zero)
<code>cmovl</code>	move if result was negative
<code>cmovle</code>	move if result was negative or zero
<code>cmovg</code>	move if result was positive
<code>cmovge</code>	result was positive or zero

* The destination operand must be a register. The source operand can be either a register or memory.

Conditional move examples

- Here is some code to compute absolute value of rax

```
mov     rbx, rax    ; save original value
neg     rax         ; negate rax
cmovl   rax, rbx    ; replace rax if negative
```

- The code below loads a number from memory, subtracts 100 and replaces the difference with 0 if the difference is negative

```
mov     rbx, 0      ; set rbx to 0
mov     rax, [x]    ; get x from memory
sub     rax, 100    ; subtract 100 from x
cmovl   rax, rbx    ; set rax to 0 if rax was negative
```

Why use a register?

- Don't use a register if a value is needed for 1 instruction
- Don't worry about it for things which execute infrequently
- Use registers instead of memory for instructions which execute enough to matter
- If you are writing a program for a class and efficiency is not part of the grade, pick the clearest way to write the code
- With so many registers, it can create opportunities for efficiency at the cost of clarity

Print to Console

```
mov  rax, 1          ; write
mov  rdi, 1          ; stdout
mov  rsi, output      ; address of first byte in output
mov  rdx, [length]    ; load length in rdx
syscall
```


Bit Operations

Bit usage

- A bit can mean one of a pair of characteristics
- True or false
- Male or female
- Bit fields can represent larger classes
 - ▶ There are 64 squares on a chess board, 6 bits could specify a position
 - ▶ The exponent field of a float can be represented using a number of bits.
 - ▶ We could use a 3 bit field to store a color from black, red, green, blue, yellow, cyan, purple and white

Bit operations

- Individual bits have values 0 and 1
- There are instructions to perform bit operations
- Using 1 as true and 0 as false
 - ▶ 1 and 1 = 1, or in C/C++, `1 && 1 = 1`
 - ▶ 1 and 0 = 0, or in C/C++, `1 && 0 = 0`
 - ▶ 1 or 0 = 1, or in C/C++, `1 || 0 = 1`
- We are interested in operations on more bits
 - ▶ `10101000b & 11110000b = 10100000b`
 - ▶ `10101000b | 00001010b = 10101010b`
- These are called “bit-wise” operations
- We will not use bit operations on single bits, though we will be able to test/set/reset individual bits

The Not operation

- C/C++ uses ! for a logical not
- C/C++ uses ~ for a bit-wise not

!0 == 1

!1 == 0

~(false) == true

~(true) == false

~10101010b == 01010101b

~0xff00 == 0x00ff

!1000000 == 0 (non-zero integer is seen as true in c/c++)

~0== ?

~1== ?

The Not operation

- C/C++ uses ! for a logical not
- C/C++ uses ~ for a bit-wise not

`!0 == 1`

`!1 == 0`

`~(false) == true`

`~(true) == false`

`~10101010b == 01010101b`

`~0xff00 == 0x00ff`

`!1000000 == 0` (non-zero integer is see as true in c/c++)

`~0 == -1`

`~1 == -2`

The Not instruction

- The `not` instruction flips all the bits of a number - one's complement
- The `not` operator does not affect any flags
- There is only a single operand which is the source and destination
- For memory operands you must include a size prefix
- The sizes are `byte`, `word`, `dword` and `qword`

```
not    rax          ; invert all bits of rax
not    dword [x]    ; invert double word at x
not    byte [x]     ; invert a byte at x
```

And operation

$\&$	0	1
0	0	0
1	0	1

- C/C++ uses `&&` for a logical and
- C/C++ uses `&` for a bit-wise and

`11001100b & 00001111b == 00001100b`

`11001100b & 11110000b == 11000000b`

`0xabcdefab & 0xff == 0xab`

`0x0123456789abcdef & 0xff00ff00ff00ff00 == 0x010045008900cd00`

- Bit-wise **and** is a bit selector

And instruction

- The **and** instruction performs a bit-wise **and**
- It has 2 operands, a destination and a source
- The source can be an immediate value, a memory location or a register
- The destination can be a register or memory
- Both destination and source cannot be in memory
- The sign flag and zero flag are set (or cleared)

And Example

- We wish to extract bits 0-3 and store them in rbx

```
mov    rax, 0x12345678
mov    rbx, rax
and    rbx, 0xf           ; rbx has the low nibble 0x8
```

- We wish to extract bits 4-7 and store them in rax

```
mov    rdx, 0             ; prepare to divide
mov    rcx, 16            ; by 16
idiv   rcx                ; rax has 0x1234567
and    rax, 0xf           ; rax has the nibble 0x7
```

Or operation

	0	1
0	0	1
1	1	1

- C/C++ uses `||` for a logical or
- C/C++ uses `|` for a bit-wise or

`11001100b | 00001111b == 11001111b`

`11001100b | 11110000b == 11111100b`

`0xabcdefab | 0xff == 0xabcdefff`

`0x0123456789abcdef | 0xff00ff00ff00ff00 == 0xff23ff67ffabffef`

- **or** is a bit setter

Or instruction

- The **or** instruction performs a bit-wise **or**
- It has 2 operands, a destination and a source
- The source can be an immediate value, a memory location or a register
- The destination can be a register or memory
- Both destination and source cannot be in memory
- The sign flag and zero flag are set (or cleared)

Or example

- Make a number odd

```
mov    rax, 0x1124
or     rax, 1           ; make the number odd
```

- Set bits 8-15.

```
mov    rax, 0x1000
or     rax, 0xff00     ; set bits 15-8
```

How would you make a number even?

Exclusive or operation

\wedge	0	1
0	0	1
1	1	0

- C/C++ uses \wedge for exclusive or

`00010001b ^ 00000001b == 00010000b`

`01010101b ^ 11111111b == 10101010b`

`01110111b ^ 00001111b == 01111000b`

`0xaaaaaaaa ^ 0xffffffff == 0x55555555`

`0x12345678 ^ 0x12345678 == 0x00000000`

- Exclusive or is a bit flipper

Exclusive or instruction

- The **xor** instruction performs a bit-wise **exclusive or**
- It has 2 operands, a destination and a source
- The source can be an immediate value, a memory location or a register
- The destination can be a register or memory
- Both destination and source cannot be in memory
- The sign flag and zero flag are set (or cleared)
- `mov rax, 0` uses 7 bytes
- `xor rax, rax` uses 3 bytes
- `xor eax, eax` uses 2 bytes

Exclusive or example

- Zero out a register.

```
mov     rax, 0x12345678
xor     eax, eax                ; set rax to 0
```

- Flip bits 0-3

```
mov     rax, 0x1234
xor     rax, 0xf                ; change to 0x123b
```

- Swap the value in two registers

```
xor rax, rbx
xor rbx, rax
xor rax, rbx
```

Shift operations

- C/C++ uses << for shift left and >> for shift right
- Shifting left introduces low order 0 bits
- Shifting right propagates the sign bit in C++ for signed integers
- Shifting right introduces 0 bits in C++ for unsigned integers
- Shifting left is like multiplying by a power of 2
- Shifting right is like dividing by a power of 2

101010b >> 3 == 101b

111111b << 2 == 11111100b

125 << 2 == 500 (125=>1111101<<2==111110100=>500)

0xabcd >> 4 == 0xabc

Shift instructions

- Shift left: `shl`
- Shift right: `shr`
- Shift arithmetic left: `sal`
- Shift arithmetic right: `sar`
- `shl` and `sal` are the same
- `shr` introduces 0 bits on the top end
- `sar` propagates the sign bit
- All the shifts use 2 operands
 - ▶ A destination register or memory
 - ▶ In immediate number of bits to shift
 - ★ Or from old 16 bit asm the `c1` register can be used
- The sign and zero flags are set (or cleared)
- The carry flag is set to the last bit shifted out

Extracting a bit field

- There are at least 2 ways to extract a bit field
- Shift right followed by an **And** operation
 - ▶ To extract bits k to m (inclusive) with $m \geq k$, shift right k bits
 - ▶ And this value with a mask of $m - k + 1$ bits all set to 1

Extracting a bit field with shift/and

Need to extract bits 9–3

1	1	0	0	0	1	1	1	1	0	0	1	0	1	1	0
15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0

Shift right 3 bits

0	0	0	1	1	0	0	0	1	1	1	0	0	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

And with 0x7f

0	0	0	0	0	0	0	0	0	1	1	1	0	0	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Extracting a bit field

- The second way
- Shift left and then right
 - ▶ Shift left until bit m is the highest bit
 - ▶ With 64 bit registers, shift left $63 - m$ bits
 - ▶ Shift right to get original bit k in position 0
 - ▶ With 64 bit registers, shift right $63 - (m - k)$ bits

Extracting a bit field with shift/shift

Need to extract bits 9–3

1	1	0	0	0	1	1	1	1	0	0	1	0	1	1	0
15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0

Shift left 6 bits

1	1	1	0	0	1	0	1	1	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Shift right 9 bits

0	0	0	0	0	0	0	0	0	1	1	1	0	0	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Rotate instructions

- The `ror` instruction rotates the bits of a register or memory location to the right
 - ▶ Values from the low end start filling in the top bits
- The `rol` instruction rotates left
 - ▶ Values from the top end of the value start filling in the low order bits
- These are 2 operand instructions like the shift instructions
- The first operand is the source to rotate (and the destination)
- The second operand is the number of bits to rotate
- The second operand is either an immediate value or `cl`
- Assuming 16 bit rotates

`1 ror 2 = 0100000000000000b`

`0xabcd ror 4 = 0xdabc`

`0x4321 rol 4 = 0x3214`

Filling a field

- There are at least 2 ways of filling in a field (with existing values)
- Use shifts and a mask.
 - ▶ Working with a 64 bit register, filling bits k to m (inclusive)
 - ▶ Prepare a mask of $m - k + 1$ bits all 1
 - ▶ Shift the new value and the mask left k bits
 - ▶ Negate the mask
 - ▶ And the old value and the mask
 - ▶ Or in the new value for the field

Filling a field 1

	We want to replace bits 6-3															
Original	1	1	0	0	0	1	1	1	1	0	0	1	0	1	0	0
	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
	with															
Value													1	1	0	1
	create mask of length $6-3+1=4$															
Mask	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1
	Shift both by $k=3$															
Value	0	0	0	0	0	0	0	0	0	1	1	0	1	0	0	0
Mask	0	0	0	0	0	0	0	0	0	1	1	1	1	0	0	0
	Negate the mask															
Mask	1	1	1	1	1	1	1	1	1	0	0	0	0	1	1	1
	And with original															
Original	1	1	0	0	0	1	1	1	1	0	0	1	0	1	0	0
Mask	1	1	1	1	1	1	1	1	1	0	0	0	0	1	1	1
	1	1	0	0	0	1	1	1	1	0	0	0	0	1	0	0
	or value with result															
	1	1	0	0	0	1	1	1	1	1	1	0	1	1	0	0

Filling a field

- Second method
- Use rotate and shift instructions and or in new value
 - ▶ Rotate the register right k bits
 - ▶ Shift the register right $m - k + 1$ bits
 - ▶ Shift the register left $m - k + 1$ bits
 - ▶ Or in the new value
 - ▶ Rotate the register left k bits

Filling a field 2

					We want to replace bits 6-3											
Original	1	1	0	0	0	1	1	1	1	0	0	1	0	1	0	0
	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Value													1	1	0	1
	Rotate original right by k=3															
	1	0	0	1	1	0	0	0	1	1	1	1	0	0	1	0
	Shift right by $m-k+1=6-3+1=4$															
	0	0	0	0	1	0	0	1	1	0	0	0	1	1	1	1
	shift left by 4															
	1	0	0	1	1	0	0	0	1	1	1	1	0	0	0	0
	or with value															
	1	0	0	1	1	0	0	0	1	1	1	1	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	1	1	0	1
	1	0	0	1	1	0	0	0	1	1	1	1	1	1	0	1
	Rotate left by k=3															
	1	1	0	0	0	1	1	1	1	1	1	0	1	1	0	0

Bit testing and setting

- It takes a few instructions to extract or set bit fields
- The same technique could be used to test or set single bits
- It can be more efficient to use special instructions operating on a single bit
 - ▶ The `bt` instruction tests a bit
 - ★ the CF flag gets set to the value of the tested bit
 - ★ we can gain access to the flag using `setc c1` (for example)
 - ▶ `bts` tests a bit and sets it
 - ★ tested bit gets set to 1
 - ▶ `btr` tests a bit and resets it
 - ★ tested bit gets set to 0
 - ▶ `btc` tests a bit and flips it
 - ★ tested bit gets complemented
- These are all 2 operand instructions
- The first operand is a register or memory location
- The second is the bit to work on, either an immediate value or a register

Bit testing and setting example

- Checking if a number is odd

```
mov rax, 101
```

```
bt rax, 0
```

```
setc dl ; 1 will be stored in dl, i.e the number is odd
```

- Setting the 7th and 33rd bit of the qword A in memory to 1

```
bts qword [A], 7
```

```
bts qword [A], 33
```

Branching and Looping

Branching and looping

- So far we have only written “straight line” code
- Conditional moves gave us an avenue for trivial if like structures.
- But we really need
 - ▶ To handle code structures like if/else. So we need both conditional and unconditional branch statements
 - ▶ We need loops

Unconditional jump

- An unconditional jump is equivalent to a `goto`
- But jumps are necessary in assembly, while high level languages could exist without `goto`
- The unconditional jump looks like
`jmp label`
- The `label` can be any label in the program's text segment
- We might think of parts of the text segment as functions
 - ▶ The computer will let you jump anywhere
 - ▶ You can try to jump to a label in the data segment, which hopefully will fail
- The assembler will generate an instruction register (`rip`) relative location to jump
 - ▶ The simplest form uses an 8 bit immediate: -128 to +127 bytes
 - ▶ The next version is 32 bits: plus or minus 2 GB
 - ▶ The short version takes up 2 bytes; the longer version 5 bytes
 - ▶ The assembler figures this out for you (Yay)

Unconditional jumps can vary

It is possible to use an unconditional jump to simulate a conditional jump.

- It is possible to jump to an address stored in a register.
- We can control the value of the register using a conditional move.

```
mov rax, a
mov rbx, b
cmovl rax, rbx ; rather jmp to b if the sign flag is set
jmp rax
```

```
a:
    .....
    .....
    jmp end
b:
    .....
    .....
end:
```


Unconditional jumps can vary

- Though it is simpler to just use a conditional jump.
- However you can construct an efficient switch statement by expanding this idea
 - ▶ You need an array of addresses and an index for the array to select which address to use for the jump

Unconditional jump used as a switch

```
segment .data
switch: dq    case0
        dq    case1
        dq    case2
i:      dq    2
segment .text
global  main                ; tell linker about main
main:   mov    rax, [i]      ; move i to rax
        jmp    [switch+rax*8] ; switch ( i )

case0:
        mov    rbx, 100     ; go here if i == 0
        jmp    end

case1:
        mov    rbx, 101     ; go here if i == 1
        jmp    end

case2:
        mov    rbx, 102     ; go here if i == 2
end:    xor    eax, eax
        ret
```

Conditional jump

- First you need to execute an instruction which sets some flags
- Then you can use a conditional jump
- The general pattern is
jCC label
- The CC means a condition code

instruction	meaning	aliases	flags
jz	jump if zero	je	ZF=1
jnz	jump if not zero	jne	ZF=0
jg	jump if $>$ zero	jnle ja	ZF=0, SF=0
jge	jump if \geq zero	jnl	SF=0
jl	jump if $<$ zero	jnge js	SF=1
jle	jump if \leq zero	jng	ZF=1 or SF=1
jc	jump if carry	jb jnae	CF=1
jnc	jump if not carry	jae jnb	CF=0

Compare operation

- It can become cumbersome to always have to perform a calculation and store the result simply to use condition jump.
- This is where the compare operation comes in handy
 - ▶ `cmp`
- `cmp` takes 2 operand.
- `cmp` subtracts the second operand from the first and sets the appropriate flags.
- But, the result is not actually stored.
- At most one operand can be an immediate value.

Simple if statement

```
if ( a < b ) {  
    temp = a;  
    a = b;  
    b = temp;  
}
```

```
    mov    rax, [a]  
    mov    rbx, [b]  
    cmp    rax, rbx  
    jge    in_order  
    mov    [a], rbx  
    mov    [b], rax  
in_order:
```

If statement with an else clause

```
if ( a < b ) {  
    max = b;  
} else {  
    max = a;  
}
```

```
    mov    rax, [a]  
    mov    rbx, [b]  
    cmp    rax, rbx  
    jnl    else  
    mov    [max], rbx  
    jmp    endif  
else:   mov    [max], rax  
endif:
```

Looping with conditional jumps

- You can construct any form of loop using conditional jumps
- We will model our code after C's loops
- `while`, `do...while` and `for`
- We will also consider `break` and `continue`
- `break` and `continue` can be avoided in C, though sometimes the result is less clear
- The same consideration applies for assembly loops as well

Sum 1 to 1000

```
sum = 0;
i = 1;
while ( i <= 100 )
{
    sum +=i;
    i++;
}
```

Now the assembler version (no optimization done to keep things simple)

Sum 1 to 1000

```
    segment .data
sum dq 0
    segment .text
    global _start
_start:
    mov rcx,1 ; i=1
while:
    cmp rcx,100
    jg ewhile
    add [sum],rcx
    inc rcx
    jmp while
ewhile:
```

Counting 1 bits in a quad-word

```
sum = 0;
i = 0;
while ( i < 64 )
{
    sum += data & 1;
    data = data >> 1;
    i++;
}
```

- There are much faster ways to do this
- But this is easy to understand and convert to assembly

Counting 1 bits in a quad-word in assembly

Assume we have the following data segment:

```
segment .data
data dq 0xfedcba9876543210
sum  dq 0
```

Counting 1 bits in a quad-word in assembly

```
segment .text
global main

main:  mov     rax, [data] ; rax holds the data
      xor     ebx, ebx    ; clear since setc will fill in bl
      xor     ecx, ecx    ; i = 0;
      xor     edx, edx    ; sum = 0;
while: cmp     rcx, 64     ; while ( i < 64 ) {
      jnl     end_while   ; requires testing on opposite
      bt      rax, 0      ; data & 1
      setc    bl          ; move result of test to bl
      add     edx, ebx    ; sum += data & 1;
      shr     rax, 1      ; data = data >> 1;
      inc     rcx         ; i++;
      jmp     while       ; end of the while loop
end_while:
      mov     [sum], rdx  ; save result in memory
      xor     eax, eax    ; return 0 from main
      ret
```

Counting 1 bits in a quad-word in assembly

To be more true to the C-code. we could replace

```
bt      rax, 0
setc    bl
add     edx, ebx
```

with

```
mov     r8, rax
and     r8, 1
add     edx, r8d
```

Branching and Looping Continued

Do-while loops

- Strict translation of a while loop uses 2 jumps

while:

 some compare

 conditional jump to "ewhile"

 jump to "while"

ewhile:

- However, a do-while only requires one jump.

do_while:

 some compare

 conditional jump to "do_while"

Do-while loops

Any **while** loop can be simulated by a **do-while** loop wrapped in an **if** statement. For example

```
while ( condition )  
{  
    statements;  
}
```

can be simulated as

```
if ( condition )  
{  
    do  
    {  
        statements;  
    } while ( condition );  
}
```


Ugly C code to search through a character array

```
//Looking for an character x other than 0. Store index in n
// data is a null terminated character array
    i = 0;
    c = data[i];
    if ( c != 0 )
    do
    {
        if ( c == x )
            break;
        i++;
        c = data[i];
    } while ( c != 0 );

    n = c == 0 ? -1 : i;
```

Assume we have the following data segment

Assume we have the following data segment

```
section .data
data    db    "hello world",0
n       dq    0
x       db    'w'
```

Assembly code to search through an array

```
        mov     bl, [x]           ; value being sought
        mov     rcx, 0            ; i = 0;
        mov     al, [data+rcx]    ; c = data[i]
        cmp     al, 0             ; if ( c != 0 ) {
        jz      end_do_while      ; skip loop for empty string
do_while:
        cmp     al, bl            ; if ( c == x ) break;
        je      found
        inc     rcx               ; i++;
        mov     al, [data+rcx]    ; c = data[i];
        cmp     al, 0             ; while ( c != 0 );
        jnz     do_while
end_do_while:
        mov     rcx, -1           ; If we get here, we failed
found:   mov     [n], rcx         ; Assign either -1 or the
                                   ; index where x was found
```

Assembly code to search through an array (Using only 64 bit registers)

```
    movzx    rbx, byte[x]           ;<----  
    mov      rcx, 0  
    movzx    rax, byte [data+rcx]  ;<----  
    cmp      rax, 0                 ;<----  
    jz       end_do_while
```

do_while:

```
    cmp      rax, rbx  
    je       found  
    inc      rcx  
    movzx    rax, byte [data+rcx]  ;<----  
    cmp      rax, 0                 ;<----  
    jnz      do_while
```

end_do_while:

```
    mov      rcx, -1
```

```
found: mov    [n], rcx
```

Counting loops

```
// assume we have 3 arrays of size n.  
// Each containing longs (quad words)
```

```
for ( i = 0; i < n; i++ )  
{  
    c[i] = a[i] + b[i];  
}
```

Counting loops

```
//assume there are 3 contiguous segments in memory
// each containing n quad words.
    mov     rdx, [n]           ; use rdx for n
    xor     ecx, ecx           ; i = 0
for:
    cmp     rcx, rdx           ; i < n
    je      end_for            ; get out if equal
    mov     rax, [a+rcx*8]     ; get a[i]
    add     rax, [b+rcx*8]     ; a[i] + b[i]
    mov     [c+rcx*8], rax     ; c[i] = a[i] + b[i];
    inc     rcx                ; i++
    jmp     for
end_for:
```

Nested loops

Consider the double summation

$$\sum_{i=1}^N \sum_{j=1}^i j \quad (1)$$

ignoring the fact that

$$\sum_{i=1}^N \sum_{j=1}^i j = \frac{N(N+1)(N+2)}{6} \quad (2)$$

We will code this double sum in assembler.

Nested loops

Assuming we have:

```
segment .data  
Sum:    dq 0  
N:      dq 5
```


Nested loops

```
    mov rbx, [N]
    mov rax, 0          ; sum=0
    mov r8,1           ; i=1
loop1:
    cmp r8,rbx
    jg eloop1          ; !(i<=N)
    mov rcx,1          ; j=1
    loop2:
        cmp rcx, r8
        jg eloop2      ; !(j<=i)
        add rax,rcx
        inc rcx
        jmp loop2
    eloop2:
        inc r8
        jmp loop1
eloop1:
    mov [Sum], rax
```

Loop instructions

- The CPU has instructions like

- ▶ `loop`
- ▶ `loope`
- ▶ `loopne`

which are designed for looping.

- They decrement `rcx` and do the branch if `rcx` is not 0
 - ▶ `loope` checks if zero flag is set as well.
 - ▶ `loopne` checks if the zero flag is not set as well
- It is faster to use `dec` and `jnz` instead
- The label must be within -128 to +127 bytes of `rip`
- Probably pointless on modern architecture. (it was fast on old architecture.)

Loop instructions

Add 5 to a sum 64 times.

```
xor rax,rax; sum=0
mov rcx, 64;
loop1:
    add rax, 5
    loop loop1
```

string (array)

- Let us first consider the simple array instruction `movsb`
- we must load the address of source data in `rsi`
- we must load the address of destination data in `rdi`
- on execution `movsb` will move the value at `rsi` to `rdi`, and increment both addresses by 1.

```
....  
mes1: db "abcdefg"  
mes2: db "1234567"  
....  
    mov rsi, mes1  
    mov rdi, mes2  
    movsb
```

`mes2` will equal "a234567" and `rsi=mes1+1` and `rdi=mes2+1`

Repeat string (array) instructions

But how is this useful?

- we utilize the string operation with the `rep` instruction.
- `rep` will repeatedly call the string operation until `rcx = 0`.

For example, let us copy an array of 1000 bytes.

```
lea rsi, [source]
lea rdi, [destination]
mov rcx, 1000
rep movsb
```

- `lea`?
- `lea rsi, [source] = mov rsi, source`

Repeat string (array) instructions

- What is if the array contains non bytes?
- We simple use the different size specifier
 - ▶ `movsw`
 - ▶ `movsd`
 - ▶ `movsq`
- Now instead of incrementing the addresses by 1,
- We will increase it by the size.
- e.g. with `movsq` we increment by 8

Repeat string (array) instructions

- Up to now we have relied on an incrementing the source and destination addresses.
- actually the address is only increased if the direction flag(DF)=0 (default)
- If DF=1, the addresses will decrement after each string instruction
- We can set the direction flag to 1 with
 std
- or 0 with
 cld

Store instruction

- The `stosb` instruction stores the byte in `al` at the address specified in `rdi` and increments `rdi`
- If the direction flag is set it decrements `rdi`
- There are also `stosw`, `stosd` and `stosq` to operate 2, 4 and 8 byte quantities

```
mov     eax, 1
mov     ecx, 1000000
lea     rdi, [destination]
rep     stosd      ; place 1000000 1's in destination
```


Scan instruction and Compare instruction

Scan: (scas)

- There are a collection of scan string instructions which scan data from the address pointed at by `rdi` and increment (or decrement) `rdi`
- They compare data against `al`, `ax`, `eax`, ...
- Used with `repe`, and will stop once data found or `rcx=0`


Compare: (cmps)

- The compare string instructions compare the data pointed at by `rdi` and `rsi`
- End once `rcx` has reached zero or a match is found.
- Used with `repe`, and will stop once match found or `rcx=0`

The Stack and Functions

The Stack

- Up until now we have mostly ignored the stack.
- Instead we have relied on the data/bss segment.
- This has been sufficient so far, but this will change.
- We are now in a position where our programs will start being complex enough that we cannot rely entirely on fixed size allocations.



The stack in a Familiar Context

- Consider the following (intentionally inefficient) recursive C++ function. Assuming we don't know that

$$fib(n) = \frac{(1 + \sqrt{5})^n - (1 - \sqrt{5})^n}{2^n \sqrt{5}} \quad (1)$$

```
// assuming fib(1)=1
int fib(int n)
{
    if(n<2)
        return 1;
    int fibM1=fib(n-1);
    int fibM2=fib(n-2);
    return fibM1+fibM2;
}
```

- If N is too large our program will experience a stack overflow.

The Stack in a Familiar Context

- But why?

The Stack in a Familiar Context

- But why?
- There are two causes (though in a sense they are the same)
 - ▶ The more obvious one is each execution of the *fib* will push at least two integers onto the stack. One for each of *FibM1*, and *FibM2*. Now the space complexity of this naive fibonacci algorithm is $O(n)$. We can try and roughly calculate *fib*(10 000 000) we need $2 * 4\text{bytes} * 10\,000\,000 = 76.29\text{mb}$ just to store the max required number of *FibM1*s, and *FibM2*s.
 - ★ Most Linux distros have a stack between *8mb* and *16mb*

The Stack in a Familiar Context

- But why?
- There are two causes (though in a sense they are the same)
 - ▶ The more obvious one is each execution of the *fib* will push at least two integers onto the stack. One for each of *FibM1*, and *FibM2*. Now the space complexity of this naive fibonacci algorithm is $O(n)$. We can try and roughly calculate *fib*(10 000 000) we need $2 * 4\text{bytes} * 10\,000\,000 = 76.29\text{mb}$ just to store the max required number of *FibM1*s, and *FibM2*s.
 - ★ Most Linux distros have a stack between *8mb* and *16mb*
- What if we implement our recursive function to somehow not push any local variable data onto the stack?
 - ▶ Even if this was possible there is one stack item we would have to push.

The Stack in a Familiar Context


- But why?
- There are two causes (though in a sense they are the same)
 - ▶ The more obvious one is each execution of the *fib* will push at least two integers onto the stack. One for each of *FibM1*, and *FibM2*. Now the space complexity of this naive fibonacci algorithm is $O(n)$. We can try and roughly calculate *fib*(10 000 000) we need $2 * 4\text{bytes} * 10\,000\,000 = 76.29\text{mb}$ just to store the max required number of *FibM1*s, and *FibM2*s.
 - ★ Most Linux distros have a stack between *8mb* and *16mb*
- What if we implement our recursive function to somehow not push any local variable data onto the stack?
 - ▶ Even if this was possible there is one stack item we would have to push.
 - ▶ The return address.

The Stack

- Now that we have seen how it can break, lets us use the stack.
- The first point to recall is that the stack starts at :0x7fffffff. Unless there is some stack randomization in place.
- The second point is that the *rsp* register stores the stack pointer.
- If nothing has been added to the stack $rsp = 0x7fffffff$.
- We can interact with the stack in a number of ways, simplest of which are the `push` and `pop` instructions.

The Stack

- The `push` instruction decrements the `rsp` register and stores the value being pushed at this address
- The `pop` instruction places the value at the top of the stack into its operand and increments `rsp`
- With the x86-64 instructions you should push and pop 8 bytes at a time
 - ▶ It is also possible to push and pop 2 bytes (word) at a time
 - ▶ Direct 4 byte push and pops is not enabled in 64-bit mode



A Stack Example

Example with start of the stack at `0x7fffffff`

```
;A
mov  rax,74
push rax
;B
inc  rax
push rax
;C
inc  rax
push rax
;D
pop  rax
;E
pop  rax
;F
```

0x7fffffff0	x	74	74	74	74	74
0x7ffffffe8	x	x	75	75	75	x
0x7ffffffe0	x	x	x	76	x	x
.	x	x	x	x	x	x
.	x	x	x	x	x	x
.	x	x	x	x	x	x
	A	B	C	D	E	F

In GDB

- You can use `x/1dg $rsp` to test the stack content.
- You can use `p/x $rsp` to test the stack pointer itself.
- Just remember that stack randomization effects the initial stack pointer.

The Stack

- Stack space is often reserved for local variables by subtracting the size needed from the stack pointer (`rsp`).
- Then an offset is used to refer to the variables.

```
sub rsp,      16 ;subtract 16 bytes
mov qword [rsp+8], 123 ;set our first qword variable to 123
mov qword [rsp],   24 ;set our second qword variable to 24
```

The stack

0x7fffffff				
0x7fffffff0	x		123	123
0x7fffffff08	x			24
0x7fffffff0e	x	x	x	x
.	x	x	x	x
.	x	x	x	x
.	x	x	x	x
	x	x	x	x
	sub rsp, 16 mov [rsp+8],123 mov [rsp], 24			

The Stack

- Remembering offsets for a large number of variables can become a burden. Instead use the `equ` pseudo-op.

```
first    equ 8
second   equ 0
mov qword [rsp+first],123
mov qword [rsp+second],24
```

The Stack

- Remembering offsets for a large number of variables can become a burden. Instead use the `equ` pseudo-op.

```
first    equ 8
second   equ 0
mov qword [rsp+first],123
mov qword [rsp+second],24
```

- How do we delete the variables after use?

The Stack

- Remembering offsets for a large number of variables can become a burden. Instead use the `equ` pseudo-op.

```
first    equ 8
second   equ 0
mov qword [rsp+first],123
mov qword [rsp+second],24
```

- How do we delete the variables after use?
 - ▶ We just move the stack pointer back.
`add rsp,16`

Functions

- Assuming we have loaded our parameters (will be explained shortly)
- You can call a function using

```
call    my_function
```

- `my_function` should be an appropriate address/label in the code segment
- The function's return value will be in `rax` or `xmm0`
- The effect of a function call is much like

```
push    next_instruction  
jmp     my_function
```

```
next_instruction:
```

The Return Instruction

- You can return to the location a function was called from using `ret`
- The effect of the return instruction (`ret`) is to pop an address off the stack and branch to it
- We could get much the same effect using

```
pop    rdi  
jmp    rdi
```

Function Parameters

- How do we get our parameters to our function?
- On 32 bit Linux all parameters were pushed onto the stack

Function Parameters

- How do we get our parameters to our function?
- On 32 bit Linux all parameters were pushed onto the stack
- On x86-64 there are 8 more registers, so some parameters are passed in registers.

Function Parameters

- How do we get our parameters to our function?
- On 32 bit Linux all parameters were pushed onto the stack
- On x86-64 there are 8 more registers, so some parameters are passed in registers.
- Linux and Mac OS/X pass integer and address parameters 1 through 6 in `rdi`, `rsi`, `rdx`, `rcx`, `r8` and `r9`
 - ▶ The remaining integer and address parameters are pushed onto the stack
 - ▶ The first 6 floating point parameters are passed in registers `xmm0` - `xmm5`

Function Parameters

- How do we get our parameters to our function?
- On 32 bit Linux all parameters were pushed onto the stack
- On x86-64 there are 8 more registers, so some parameters are passed in registers.
- Linux and Mac OS/X pass integer and address parameters 1 through 6 in `rdi`, `rsi`, `rdx`, `rcx`, `r8` and `r9`
 - ▶ The remaining integer and address parameters are pushed onto the stack
 - ▶ The first 6 floating point parameters are passed in registers `xmm0` - `xmm5`
- Windows uses registers `rcx`, `rdx`, `r8` and `r9` for the first 4 integer and address parameters
 - ▶ The remaining integer and address parameters are pushed onto the stack. (there is also a set amount of stack padding)
 - ▶ Windows uses `xmm0` - `xmm3`

Function Parameters

- How do we get our parameters to our function?
- On 32 bit Linux all parameters were pushed onto the stack
- On x86-64 there are 8 more registers, so some parameters are passed in registers.
- Linux and Mac OS/X pass integer and address parameters 1 through 6 in `rdi`, `rsi`, `rdx`, `rcx`, `r8` and `r9`
 - ▶ The remaining integer and address parameters are pushed onto the stack
 - ▶ The first 6 floating point parameters are passed in registers `xmm0` - `xmm5`
- Windows uses registers `rcx`, `rdx`, `r8` and `r9` for the first 4 integer and address parameters
 - ▶ The remaining integer and address parameters are pushed onto the stack. (there is also a set amount of stack padding)
 - ▶ Windows uses `xmm0` - `xmm3`
- In all cases pushed parameters are pushed in reverse order

Function Parameters (2)

- Functions like `printf` having a variable number of parameters must place the number of floating point parameters in `rax`

Register	Usage	Preserved across function calls
<code>rax</code>	temporary register; with variable arguments passes information about the number of vector registers used; 1 st return register	No
<code>rbx</code>	callee-saved register; optionally used as base pointer	Yes
<code>rcx</code>	used to pass 4 th integer argument to functions	No
<code>rdx</code>	used to pass 3 rd argument to functions; 2 nd return register	No
<code>rsp</code>	stack pointer	Yes
<code>rbp</code>	callee-saved register; optionally used as frame pointer	Yes
<code>rsi</code>	used to pass 2 nd argument to functions	No
<code>rdi</code>	used to pass 1 st argument to functions	No
<code>r8</code>	used to pass 5 th argument to functions	No
<code>r9</code>	used to pass 6 th argument to functions	No
<code>r10</code>	temporary register, used for passing a function's static chain pointer	No
<code>r11</code>	temporary register	No
<code>r12-r15</code>	callee-saved registers	Yes

Simple Function

Simple function that returns the larger of two longs.

```
; long max(long a, long b)
```

```
max:
```

```
    mov rax, rdi    ; move parm1 to rax
```

```
    cmp rax, rsi    ; compare rax to parm2
```

```
    cmovl rax, rsi  ; if parm2 > rax then move parm 2 to rax
```

```
    ret
```

Simple Function

Calling the simple function

```
mov rdi, 123 ; load parm1  
mov rsi, 742 ; load parm2  
call max
```

Real World Functions

- Both Linux and Windows require the maintenance of the stack on 16 byte boundaries during the main part of functions.

Real World Functions

- Both Linux and Windows require the maintenance of the stack on 16 byte boundaries during the main part of functions.
 - ▶ The reason behind this requirement is to make it possible for local variables (on the stack) to be on 16 byte boundaries,
 - ▶ a requirement for some SSE(Streaming SIMD Extensions) and AVX(Advanced Vector Extension) instructions.

Real World Functions

- Both Linux and Windows require the maintenance of the stack on 16 byte boundaries during the main part of functions.
 - ▶ The reason behind this requirement is to make it possible for local variables (on the stack) to be on 16 byte boundaries,
 - ▶ a requirement for some SSE(Streaming SIMD Extensions) and AVX(Advanced Vector Extension) instructions.
- Conforming functions generally start with “push rbp” which re-establishes the 16 byte bounding temporarily **botched by the function call**
 - ▶ Remember, the call operation pushes a 8 byte value onto the stack (the return address)

Real World Functions

- Both Linux and Windows require the maintenance of the stack on 16 byte boundaries during the main part of functions.
 - ▶ The reason behind this requirement is to make it possible for local variables (on the stack) to be on 16 byte boundaries,
 - ▶ a requirement for some SSE(Streaming SIMD Extensions) and AVX(Advanced Vector Extension) instructions.
- Conforming functions generally start with “push rbp” which re-establishes the 16 byte bounding temporarily **botched by the function call**
 - ▶ Remember, the call operation pushes a 8 byte value onto the stack (the return address)
- Following that, conforming functions subtract multiples of 16 from rsp to allocate stack space or push pairs of 8 byte values
 - ▶ Even if this means over allocation

Hello world, at last

```
        section .data
msg:     db      "Hello World!",0x0a,0

        section .text
global  main
extern  printf

main:
    push    rbp
    mov     rbp, rsp      ; will explain shortly
    mov     rdi, msg      ; parameter 1 for printf
    mov     rax, 0        ; 0 floating point parameters
    call    printf
    mov     rax, 0        ; return 0
    mov     rsp, rbp      ; will explain shortly (NIB)
    pop     rbp
    ret
```

Stack frames

- Stack frames are used by the gdb debugger to trace backwards through the stack to inspect calls made in a process
- If we start and end each function like:

```
push    rbp
mov     rbp, rsp
```

```
...
mov     rsp, rbp
pop     rbp
ret
```

- We are in effect constructing a link list of all of the stack frames.
- All non-leaf functions must have the stack frame set up and destruction to conform to the ABI and be properly c/c++ compatible.

Stack frames

	In base function:	In Function L1:	In Function L2:	In Function L3:	In Function L4:
rbp	0x0	StackPointer(0)	StackPointer(1)	StackPointer(2)	StackPointer(3)
rsp	StackPointer(0)	StackPointer(1)	StackPointer(2)	StackPointer(3)	StackPointer(4)
0x7fffffff0	x	Ret Adr to Base	Ret Adr to Base	Ret Adr to Base	Ret Adr to Base
0x7fffffff8	x	0x0	0x0	0x0	0x0
0x7fffffffe0	x	x	Ret Adr to L1	Ret Adr to L1	Ret Adr to L1
0x7fffffffd8	x	x	StackPointer(0)	StackPointer(0)	StackPointer(0)
0x7fffffffd0	x	x	x	Ret Adr to L2	Ret Adr to L2
0x7fffffffc8	x	x	x	StackPointer(1)	StackPointer(1)
0x7fffffffc0	x	x	x	x	Ret Adr to L3
0x7fffffffb8	x	x	x	x	StackPointer(2)
0x7fffffffb0	x	x	x	x	x

*assuming no local variables are stored on the stack. If the local variables were stored where would they be?

Stack frames

- If you require space for local variables you simple need to subtract an amount from the stack pointer `rsp`.

Stack frames

- If you require space for local variables you simple need to subtract an amount from the stack pointer `rsp`.
 - ▶ The subtraction should always maintain the 16 byte boundary.
 - ▶ For example, say we wish to have a single quadword as a local variable.

Stack frames

- If you require space for local variables you simple need to subtract an amount from the stack pointer `rsp`.
 - ▶ The subtraction should always maintain the 16 byte boundary.
 - ▶ For example, say we wish to have a single quadword as a local variable.
 - ★ We only **have** to subtract 8 bytes.
 - ★ But we **should** maintain the 16 byte boundary.
- ```
push rbp
mov rbp, rsp
sub rsp,16
```

# Stack frames

- If you require space for local variables you simple need to subtract an amount from the stack pointer `rsp`.
  - ▶ The subtraction should always maintain the 16 byte boundary.
  - ▶ For example, say we wish to have a single quadword as a local variable.
    - ★ We only **have** to subtract 8 bytes.
    - ★ But we **should** maintain the 16 byte boundary.

```
push rbp
mov rbp, rsp
sub rsp,16
```

- if we allocate local variables we **must** use

```
mov rsp, rbp
pop rbp
ret
```

Just popping will not work.

# Stack frames

- If you prefer you can utilize

`leave`

`ret`

instead of

`mov rsp, rbp`

`pop rbp`

`ret`

## print max example

```
main:
 push rbp
 mov rbp, rsp
; print_max (100, 200);
 mov rdi, 100 ; first parameter
 mov rsi, 200 ; second parameter
 call print_max
 mov rax, 0 ; to return 0
 leave
 ret
```

## print max example

```
; void print_max (long a, long b)
; {
a equ 0
b equ 8
max equ 16
print_max:
 push rbp
 mov rbp, rsp
 sub rsp, 32 ; leave space for a, b and max
 mov [rsp+a], rdi ; save a
 mov [rsp+b], rsi ; save b
 mov [rsp+max], rdi ; max = a;
 cmp rsi, rdi ; if (b > max) max = b
 jng skip
 mov [rsp+max], rsi
```



## print max example

```
skip:
segment .data
 fmt db "max(%ld,%ld) = %ld",0xa,0
segment .text
 mov rdi, fmt ; address of format string
 mov rsi, [rsp+a] ; first %ld
 mov rdx, [rsp+b] ; second %ld
 mov rcx, [rsp+max] ; third %ld
 mov rax, 0 ; zero floating point param
 call printf
 leave
 ret
```

# Accessing Local Variables

- Can you see the potential issue with the way we are accessing our local variables?

# Accessing Local Variables

- Can you see the potential issue with the way we are accessing our local variables?
  - ▶ If anything is pushed onto the stack our offsets become incorrect.

# Accessing Local Variables

- Can you see the potential issue with the way we are accessing our local variables?
  - ▶ If anything is pushed onto the stack our offsets become incorrect.
- Solution?

# Accessing Local Variables

- Can you see the potential issue with the way we are accessing our local variables?
  - ▶ If anything is pushed onto the stack our offsets become incorrect.
- Solution?
  - ▶ Simply use `rbp` as the base address. Just remember that since `rbp` holds the value of `rsp` before the subtraction you need to use subtract from `rbp` and not add (think about what the offset values should be).
  - ▶ or just don't use push and pop other than for the stack frame setup and destruction

# Recursive Functions

- Recursive algorithms serve as a good example for why we need stack based storage.
- Often times we can get away with utilizing registers that are preserved across function calls.
- However consider the case where we have a recursive algorithm. On the first level we decide to use `r15` to store a value. But now on the second level `r15` is already in use...

# Recursive Functions

- Consider again the following recursive Fibonacci function

```
// assuming fib(1)=1
long fib(long n)
{
 if(n<2)
 return 1;
 return fib(n-1)+fib(n-2);
}
```

# Recursive Functions

```
• fib: push rbp
 mov rbp, rsp
 sub rsp, 16
N equ 0
nM1 equ 8
 mov rax, 1 ;base case return value
 cmp rdi, 2
 jl .end ;first parameter<2 (base case)
 dec rdi ;recall rdi is the first parameter (n)
 mov [rsp+N], rdi ;save N-1
 call fib
 mov [rsp+nM1], rax
 mov rdi, [rsp+N] ;load N-1
 dec rdi
 call fib
 add rax, [rsp+nM1]
 .end
 leave
 ret
```



# Function Implementation: Correct Practice

- If a function is a **non-leaf** function you **must** set up and destroy a stack frame.
- If you utilize a register in your function that should be preserved across function calls (like r15) you **must** restore it to its original value.
- Example

```
sub rsp, 16
mov [rsp], r15
mov [rsp+8], r14
....
....
mov r14, [rsp+8]
mov r15, [rsp]
add rsp, 16
```

# Function Implementation: Correct Practice

- or if you make use of no other stack based memory.

```
push r15
```

```
push r14
```

```
....
```

```
....
```

```
pop r14
```

```
pop r15
```

## Floating Point Instructions

# Floating point instructions

- PC floating point operations were once done in a separate chip - 8087
  - ▶ This chip managed a stack of 80 bit floating point values.
  - ▶ The stack and instructions still exist, but are largely ignored.
  - ▶ In the absence of an 8087 chip, floating point operation were emulated in software.

# Floating point instructions

- PC floating point operations were once done in a separate chip - 8087
  - ▶ This chip managed a stack of 80 bit floating point values.
  - ▶ The stack and instructions still exist, but are largely ignored.
  - ▶ In the absence of an 8087 chip, floating point operation were emulated in software.
- x86-64 CPUs have 16 floating point registers (128 or 256 bits)
  - ▶ These registers can be used for single data instructions or single instruction multiple data instructions (SIMD)

# Floating point instructions

- PC floating point operations were once done in a separate chip - 8087
  - ▶ This chip managed a stack of 80 bit floating point values.
  - ▶ The stack and instructions still exist, but are largely ignored.
  - ▶ In the absence of an 8087 chip, floating point operation were emulated in software.
- x86-64 CPUs have 16 floating point registers (128 or 256 bits)
  - ▶ These registers can be used for single data instructions or single instruction multiple data instructions (SIMD)
- We will focus on these newer registers
  - ▶ The older instructions tended to start with the letter “f” and referenced the stack using register names like ST0
  - ▶ The newer instructions reference using registers with names like “xmm0”, and “ymm0”. (zmm0 in new CPUs as well)

# Floating point instructions

- There are 16 floating point registers.
- `ymm0` to `ymm15` (AVX registers)
- Each one is 256 bits.
- The lower half (128 bits) of `ymm-` is referred to as `xmm-`
- `xmm0` to `xmm15` (SSE registers)
- The full 256 bit register are available from the Core i series.

# Floating point instructions

- There are 16 floating point registers.
- `ymm0` to `ymm15` (AVX registers)
- Each one is 256 bits.
- The lower half (128 bits) of `ymm-` is referred to as `xmm-`
- `xmm0` to `xmm15` (SSE registers)
- The full 256 bit register are available from the Core i series.
- We will mainly discuss `xmm` registers, all operations are the same for `ymm` registers, we just append a `v` in front of the instruction.



# Moving scalars to or from floating point registers

## Moving floating point numbers

- The two instructions available are `movss` and `movsd`
- `movss` moves a single 32 bit floating point value to or from an `xmm` register (float/single)
- `movsd` moves a single 64 bit floating point value (double)

# Moving scalars to or from floating point registers

## Moving floating point numbers

- The two instructions available are `movss` and `movsd`
- `movss` moves a single 32 bit floating point value to or from an `xmm` register (float/single)
- `movsd` moves a single 64 bit floating point value (double)
- It should be noted that there is no implicit data conversion - unlike the old instructions which converted floating point data to an 80 bit internal format

# Moving scalars to or from floating point registers

- The instructions follow the standard pattern of having at most one memory address

```
segment .data
x: dd 12.35 ; float/single
y: dq 14.36 ; double

movss xmm0, [x] ; move the float value at x into xmm0
movsd [y], xmm1 ; move double value from xmm1 to y
movss xmm2, xmm0 ; move from xmm0 to xmm2
```

# Moving packed data

- The XMM registers are 128 bits
  - ▶ They can hold 4 floats or 2 doubles (or integers of various sizes)
- The YMM registers are 256 bits
  - ▶ They can hold 8 floats or 4 doubles (or integers of various sizes)

# Moving packed data

- The XMM registers are 128 bits
  - ▶ They can hold 4 floats or 2 doubles (or integers of various sizes)
- The YMM registers are 256 bits
  - ▶ They can hold 8 floats or 4 doubles (or integers of various sizes)
- But how do we load them?
  - ▶ There are two types of packed move instructions available.
  - ▶ An **aligned** version and an **unaligned** version.
  - ▶ **aligned** move requires the data to be of a 16 byte boundary, but is faster in general.
  - ▶ **unaligned** move is slower, though more so on older CPUs.

# Moving packed data

- The XMM registers are 128 bits
  - ▶ They can hold 4 floats or 2 doubles (or integers of various sizes)
- The YMM registers are 256 bits
  - ▶ They can hold 8 floats or 4 doubles (or integers of various sizes)
- But how do we load them?
  - ▶ There are two types of packed move instructions available.
  - ▶ An **aligned** version and an **unaligned** version.
  - ▶ **aligned** move requires the data to be of a 16 byte boundary, but is faster in general.
  - ▶ **unaligned** move is slower, though more so on older CPUs.
- If you try to use an aligned move on unaligned data you will get a segmentation fault.

# Moving packed data

Actual packed move instructions:

- `movaps` moves 4 floats to/from a memory address aligned at a 16 byte boundary
- `movups` does the same task with unaligned memory addresses
- `movapd` moves 2 doubles to/from a memory address aligned at a 16 byte boundary
- `movupd` does the same task with unaligned memory addresses

```
segment .data
x: dd 12.3, 9.3, 123.2, 0.1
a: dq 0, 0
....
movups xmm0, [x] ; move 4 floats to xmm0
movupd [a], xmm15 ; move 2 doubles to a
```

# Moving packed data

If you wish to use the aligned move:

```
 segment .data
align 16
x: dd 12.3, 9.3, 123.2, 0.1
....
 movaps xmm0, [x] ; move 4 floats to xmm0
```



# Floating point addition

- `addss` adds a scalar float (single precision) to another
- `addsd` adds a scalar double to another
- `addps` adds 4 floats to 4 floats - pairwise addition
- `addpd` adds 2 doubles to 2 doubles
- There are 2 operands: destination and source
- The source can be memory or an XMM register
- The destination must be an XMM register
- Flags are unaffected

```
movss xmm0, [a] ; load a
addss xmm0, [b] ; add b to a
movss [c], xmm0 ; store sum in c
```

And

```
movapd xmm0, [a] ; load 2 doubles from a
addpd xmm0, [b] ; add a[0]+b[0] and a[1]+b[1]
movapd [c], xmm0 ; store 2 sums in c
```

# Floating point subtraction

- `subss` subtracts the source float from the destination
- `subsd` subtracts the source double from the destination
- `subps` subtracts 4 floats from 4 floats
- `subpd` subtracts 2 doubles from 2 doubles

```
movss xmm0, [a] ; load a
subss xmm0, [b] ; add b from a
movss [c], xmm0 ; store a-b in c
```

And

```
movapd xmm0, [a] ; load 2 doubles from a
subpd xmm0, [b] ; add a[0]-b[0] and a[1]-b[1]
movapd [c], xmm0 ; store 2 differences in c
```

# Basic floating point instructions

| instruction | effect                 |
|-------------|------------------------|
| addsd       | add scalar double      |
| addss       | add scalar float       |
| addpd       | add packed double      |
| addps       | add packed float       |
| subsd       | subtract scalar double |
| subss       | subtract scalar float  |
| subpd       | subtract packed double |
| subps       | subtract packed float  |
| mulsd       | multiply scalar double |
| mulss       | multiply scalar float  |
| mulpd       | multiply packed double |
| mulps       | multiply packed float  |
| divsd       | divide scalar double   |
| divss       | divide scalar float    |
| divpd       | divide packed double   |
| divps       | divide packed float    |

## Conversion to a different length floating point

- `cvtss2sd` converts a scalar single (float) to a scalar double
- `cvtps2pd` converts 2 packed floats to 2 packed doubles
- `cvtss2sd` converts a scalar double to a scalar float
- `cvtpd2ps` converts 2 packed doubles to 2 packed floats

```
cvtss2sd xmm0, [a] ; get a into xmm0 as a double
addsd xmm0, [b] ; add a double to a
cvtss2sd xmm0, xmm0 ; convert to float
movss [c], xmm0
```

# Converting floating point to/from integer

- `cvtss2si` converts a float to a double word or quad word integer by rounding
- `cvtsd2si` converts a float to a double word or quad word integer by rounding
- `cvtss2si` and `cvtsd2si` convert by truncation
- `cvtss2si` converts an integer to a float in an XMM register
- `cvtsd2si` converts an integer to a double in an XMM register
- When converting integers from memory a size qualifier is needed

```
cvtss2si eax, xmm0 ; convert to dword integer
cvtss2si rax, xmm0 ; convert to qword integer
cvtsd2si xmm0, rax ; convert qword to double
cvtsd2si xmm0, dword [x] ; convert dword integer
```

# Unordered versus ordered comparisons

- In the IEEE-754 floating point standard there are two types of NaNs (not a number)
- QNaN or SNaN
  - ▶ QNaN means “quiet, not a number”
  - ▶ SNaN means “signalling, not a number”
  - ▶ Both have all exponent field bits set to 1
  - ▶ QNaN has its top fraction bit equal to 1

# Unordered versus ordered comparisons

- Floating point comparisons can cause exceptions
- Ordered comparisons cause exceptions on QNaN or SNaN
- An unordered comparison causes exceptions only for SNaN
- gcc uses unordered comparisons
- If it's good enough for gcc, it's good enough for us
- `ucomiss` compares floats
- `ucomisd` compares doubles
- The first operand **must** be an XMM register
- They set the zero flag, parity flag and carry flags

```
movss xmm0, [a]
mulss xmm0, [b]
ucomiss xmm0, [c]
jbe less_eq ; jmp if a*b <= c
```

# Conditional floating point jumps

| instruction | meaning                | aliases | flags        |
|-------------|------------------------|---------|--------------|
| jb          | jump if below          | jc jnae | CF=1         |
| jbe         | jump if below or equal | jna     | ZF=1 or CF=1 |
| ja          | jump if above          | jnbe    | ZF=0 or CF=0 |
| jae         | jump if above or equal | jnc jnb | CF=0         |
| je          | jump if equal          | jz      | ZF=1         |
| jne         | jump if not equal      | jnz     | ZF=0         |

c= carry flag set

z= zero flag set



# Mathematical functions

- 8087 had sine, cosine, arctangent and more
- The newer instructions omit these operations on XMM registers
- Instead you are supposed to use efficient library functions
- There are instructions for
  - ▶ Minimum
  - ▶ Maximum
  - ▶ Rounding
  - ▶ Square root
  - ▶ Reciprocal of square root

# Minimum and maximum

- `minss` and `maxss` compute minimum or maximum of scalar floats
- `minsd` and `maxsd` compute minimum or maximum of scalar doubles
- The destination operand must be an XMM register
- The source can be an XMM register or memory
- `minps` and `maxps` compute minimum or maximum of packed floats
- `minpd` and `maxpd` compute minimum or maximum of packed doubles
- `minps xmm0, xmm1` computes 4 minimums and places them in `xmm0`

```
movss xmm0, [x] ; move x into xmm0
maxss xmm0, [y] ; xmm0 has max(x,y)
movapd xmm0, [a] ; move a[0] and a[1] into xmm0
minpd xmm0, [b] ; xmm0[0] has min(a[0],b[0])
 ; xmm0[1] has min(a[1],b[1])
```

# Rounding

- `roundss` rounds 1 float
- `roundps` rounds 4 floats
- `roundsd` rounds 1 double
- `roundpd` rounds 2 doubles
- The first operand is an XMM destination register
- The second is the source in an XMM register or memory
- The third operand is a rounding mode

| mode | meaning                            |
|------|------------------------------------|
| 0    | round, giving ties to even numbers |
| 1    | round down                         |
| 2    | round up                           |
| 3    | round toward 0 (truncate)          |

# Square roots

- `sqrtps` computes 1 float square root
- `sqrtps` computes 4 float square roots
- `sqrtsd` computes 1 double square root
- `sqrtpd` computes 2 double square roots
- The first operand is an XMM destination register
- The second is the source in an XMM register or memory

## Distance in 3D

$$d = \sqrt{((x_1 - x_2)^2 + (y_1 - y_2)^2 + (z_1 - z_2)^2)}$$

distance3d:

```
movss xmm0, [rdi] ; x from first point
subss xmm0, [rsi] ; subtract x from second point
mulss xmm0, xmm0 ; (x1-x2)^2
movss xmm1, [rdi+4] ; y from first point
subss xmm1, [rsi+4] ; subtract y from second point
mulss xmm1, xmm1 ; (y1-y2)^2
movss xmm2, [rdi+8] ; z from first point
subss xmm2, [rsi+8] ; subtract z from second point
mulss xmm2, xmm2 ; (z1-z2)^2
addss xmm0, xmm1 ; add x and y parts
addss xmm0, xmm2 ; add z part
sqrtss xmm0, xmm0
ret
```

# Dot product in 3D

$$d = x_1x_2 + y_1y_2 + z_1z_2$$

dot\_product:

```
 movss xmm0, [rdi]
 mulss xmm0, [rsi]
 movss xmm1, [rdi+4]
 mulss xmm1, [rsi+4]
 addss xmm0, xmm1
 movss xmm2, [rdi+8]
 mulss xmm2, [rsi+8]
 addss xmm0, xmm2
 ret
```

# Arrays

# Arrays

- An array is a **contiguous** collection of memory cells of a specific type
- The start address of an array is the address of the first element
  - ▶ The start address is associated with the label given before a data definition in the data segment or a data reservation in the bss segment.
  - ▶ Unless the array is in allocated memory.
- The first index of an array in C/C++ and assembly is 0
- Some high level languages use different or user-selectable starting indices for arrays
  - ▶ Fortran defaults to 1



# Array address computation

- Array elements all have the same size: 1, 2, 4 and 8 are common
  - ▶ If I have a pointer in my struct. The heap allocated memory does not explicitly contribute to the size of the object.
  - ▶ The pointer itself contributes to the size of the object not the heap allocated memory
  - ▶ A pointer on 64bit linux. Is 64 bits.
- Suppose an array has elements of size 4 and starts at address 0x10000
  - ▶ The first element (at index 0) is at 0x10000
  - ▶ The second element (at index 1) is at 0x10004
  - ▶ The third element (at index 2) is at 0x10008
  - ▶ Element number  $k$  is at address  $0x10000 + k*4$

# General pattern for memory references

|               |                                                                                                   |
|---------------|---------------------------------------------------------------------------------------------------|
| [label]       | the value contained at label                                                                      |
| [label+ind]   | the value contained at the memory address obtained by adding the label and index register         |
| [label+2*ind] | the value contained at the memory address obtained by adding the label and index register times 2 |
| [label+4*ind] | the value contained at the memory address obtained by adding the label and index register times 4 |
| [label+8*ind] | the value contained at the memory address obtained by adding the label and index register times 8 |

# General pattern for memory references

Consider:

```
segment .data
c: dq 4,1,5,2,7,8
```

Then,

```
mov rax, [c];
```

moves 4 into rax. And

```
mov rcx, 2
mov rax, [c+8*rcx];
```

moves 5 into rax.

What would the following misguided move load?

```
mov rcx, 1
mov rax, [c+4*rcx];
```

# General pattern for memory references

|                                |                                                                                                                 |
|--------------------------------|-----------------------------------------------------------------------------------------------------------------|
| <code>[reg]</code>             | the value contained at the memory address in the register                                                       |
| <code>[reg+k*ind]</code>       | the value contained at the memory address obtained by adding the register and index register times k            |
| <code>[label+reg+k*ind]</code> | the value contained at the memory address obtained by adding the label, the register and index register times k |
| <code>[n+reg+k*ind]</code>     | the value contained at the memory address obtained by adding n, the register and index register times k         |

# General pattern for memory references

Consider:

```
segment .data
a: dq 123
c: dq 4,1,5,2,7,8
```

Then

```
lea rcx, [c] ; or mov rcx, c
mov rax, [rcx]
```

will load 4 into rax. And

```
lea rcx, [a]
mov rdi, 4
mov rax, [8 + rcx + 8*rdi]
```

# General pattern for memory references

Consider:

```
segment .data
a: dq 123
c: dq 4,1,5,2,7,8
```

Then

```
lea rcx, [c] ; or mov rcx, c
mov rax, [rcx]
```

will load 4 into rax. And

```
lea rcx, [a]
mov rdi, 4
mov rax, [8 + rcx + 8*rdi]
```

- will load 7 into rax.

# Memory references

- For items in the data and bss segments we can use a label
- For arrays passed into functions the address is passed in a register
- Soon we will be allocating memory using `malloc`
  - ▶ This address will typically be stored in memory
  - ▶ Later to use the data, we must load the address from memory into a register
  - ▶ Then we can use a register form of memory reference
- The use of a number or a label is equivalent to the computer
  - ▶ Both use the same instruction and place the number or label value into the same field of the instruction
  - ▶ Using multipliers of 2, 4 or 8 are essentially “free” with index registers

## Copy dword array example

- In the function below the first parameter is the address of the first dword of a destination array (rdi)
- The second parameter is the address of the source array (rsi)
- The third parameter is the number of dwords to copy (rdx)
- It would generally be faster to use “rep movsd”

copy\_array:

```
 xor ecx, ecx ; index=0
more: mov eax, [rsi+4*rcx] ; move src[index] to temp
 mov [rdi+4*rcx], eax ; move to dst[index]
 inc rcx ; ++index
 cmp rcx, rdx
 jne more
 xor eax, eax
 ret
 ; if rdx=0 bad things happen
```



# Allocating arrays

If we wish to directly allocate heap storage in assembler we have two options.

- We can make use of the `brk` and `sbrk` system calls which allow us a means of altering the heap boundary.
- Or the more modern approach using `mmap`.

In this course we will however make use of the C `malloc` function.

- If `malloc` is not fast enough, your time would be better served rewriting a version of `malloc` for your purposes (maybe in ASM) rather than using the system calls directly all the time.
- A nice guide can be found at <https://moss.cs.iit.edu/cs351/slides/slides-malloc.pdf> (this is only if you are interested)

# Allocating arrays

- We will allocate arrays using the C `malloc` function

```
void *malloc (long size);
```

- The parameter to `malloc` is the number of bytes to allocate
- `malloc` returns the address of the array or 0
- Data allocated should be freed

```
void free (void *ptr);
```

# Code to allocate an array

- The code below allocates an array of 1 billion bytes
- It saves the pointer to the new array in memory location named `pointer`

```
extern malloc
...
mov rdi, 1000000000
call malloc
mov [pointer], rax
```

# Advantages for using allocated arrays

- The array will be the right size
- There are size limits of about 2 GB in the data and bss segments
- The assembler phase is very slow with large arrays and the program is large
- Assembling a program with a 2 GB array in the data segment took about 100 seconds
- The executable was over 2 GB
- Using `malloc` the program assembles in less than 1 second and the executable as about 10 KB

# Processing arrays

- We present an application which creates an array
- Fills the array with random data by calling `random`
- Prints the array if the size is small
- Determines the minimum value in the array.
- Only the helper functions will be discussed in the lecture.

# Creating an array

- This function allocates an array of **double words**
- The number of **double words** is the only parameter
- Note the use of a stack frame to avoid any problems of stack misalignment

```
; array = create (size);
```

```
create:
```

```
 push rbp
 mov rbp, rsp
 imul rdi, 4
 call malloc
 leave
 ret
```

# Filling the array with random numbers

```
fill: ; void fill(int* array,long size) \\ assumes size>=1
.array equ 0
.size equ 8
.i equ 16
 push rbp
 mov rbp, rsp
 sub rsp, 32
 mov [rsp+.array], rdi
 mov [rsp+.size], rsi
 mov rcx, 0
.more mov [rsp+.i], rcx
 call rand ;rand returns an integer (int=32 bits)
 mov rcx, [rsp+.i]
 mov rdi, [rsp+.array]
 mov [rdi+rcx*4], eax
 inc rcx
 cmp rcx, [rsp+.size]
 jl .more
 leave
 ret
```

# Printing the array

```
; void print (int* array, long size);
print:
.array equ 0
.size equ 8
.i equ 16
push rbp
mov rbp, rsp
sub rsp, 32
mov [rsp+.array], rdi
mov [rsp+.size], rsi
mov rcx, 0
mov [rsp+.i], rcx
```



# Printing the array

```
 segment .data
.format:
 db "%10d",0x0a,0
 segment .text
.more
 lea rdi, [.format]
 mov rdx, [rsp+.array]
 mov rcx, [rsp+.i]
 mov esi, [rdx+rcx*4]
 mov rax, 0
 call printf
 mov rcx, [rsp+.i]
 inc rcx
 mov [rsp+.i], rcx
 cmp rcx, [rsp+.size]
 jl .more
 leave
 ret
```

## Finding the minimum value in the array

- This function calls no other function
- There is no need for a stack frame
  - ▶ but there is no real harm in having one)
- A conditional move is faster than branching

```
; x = min (a, size); int min(int* array, long size)
; assumes size>=1
min:
 mov eax, [rdi] ; start with a[0]
 mov rcx, 1
.more mov r8d, [rdi+rcx*4] ; get a[i]
 cmp r8d, eax
 cmovl eax, r8d ; move if smaller
 inc rcx
 cmp rcx, rsi
 jl .more
 ret
```

# Command line parameter array

- The first argument to `main` is the number of command line parameters
- The second argument is the address of an array of character pointers, each pointing to one of the parameters
- Below is a C program illustrating the use of command line parameters

```
#include <stdio.h>
```

```
int main (int argc, char *argv[])
{
 int i;
 for (i = 0; i < argc; i++) {
 printf("%s\n", argv[i]);
 }
 return 0;
}
```

# Assembly program listing command line parameters

```
 segment .data
format db "%s",0x0a,0
 segment .text
 global main ; let the linker know about main
 extern printf ; resolve printf from libc
main: push rbp ; prepare stack frame for main
 mov rbp, rsp
 sub rsp, 16
 mov rcx, rsi ; move argv to rcx
 mov rsi, [rcx] ; get first argv string
start_loop:
 lea rdi, [format]
 mov [rsp], rcx ; save argv
 call printf
 mov rcx, [rsp] ; restore argv
 add rcx, 8 ; advance to next pointer in argv
 mov rsi, [rcx] ; get next argv string
 cmp rsi, 0
 jnz start_loop ; end with NULL pointer
end_loop:
```

## System Calls

# System calls

- A system call is how a program requests a service from an operating system's kernel.
- A user process cannot do privileged instructions
  - ▶ No direct access to a hard drive
  - ▶ No changing of CPU mapping registers
- Instead a user process makes a system call
- The system call is a part of the kernel of the operating system
  - ▶ It verifies that the user should be allowed to do the requested action and then does the action.

## 32 bit Linux system calls

- Each system call is identified by an integer defined in “/usr/include/asm/unistd\_32.h” (varies slightly by distro)
- The system call number is placed in `eax`
- Parameters are placed in registers `ebx`, `ecx`, `edx`, `esi`, `edi`, and `ebp`
- Processes use the software interrupt number `0x80` to make the system call
- Return value in `eax`

```
 segment .data
hello: db "Hello world!",0x0a
 segment .text
...
mov eax, 4 ; syscall 4 is write
mov ebx, 1 ; file descriptor
lea ecx, [hello] ; array to write
mov edx, 13 ; write 13 bytes
int 0x80
```

# 64 bit Linux system calls

- System call number defined in  
“/usr/include/asm/unistd\_64.h” (varies slightly by distro)
- System call number is placed in rax
- Parameters rdi, rsi, rdx, r10, r8 and r9.
- Process uses syscall instruction
- Return value in rax.

```
segment .data
hello: db "Hello world!",0x0a
segment .text
global _start
_start: mov eax, 1 ; syscall 1 is write
 mov edi, 1 ; file descriptor
 lea rsi, [hello] ; array to write
 mov edx, 13 ; write 13 bytes
 syscall
 mov eax, 60 ; syscall 60 is exit
 xor edi, edi ; exit(0)
 syscall
```



# C wrapper functions

- Every system call is available through a C “wrapper function”
- A wrapper function might do very little other than shuffle registers
- Some wrappers offer a little extra convenience
- Wrapper functions are described in section 2 of the on-line manual
  - ▶ Use “man 2 write” to learn about the write system call

```
segment .data
msg: db "Hello World!",0x0a ; String to print
len: equ $-msg ; Length of the string
segment .text
global main
extern write, exit

main:
 mov edi, 1 ; Arg 1 is the fd
 lea rsi, [msg] ; Arg 2 is the array
 mov edx, len ; Arg 3 is the length
 call write
 xor edi, edi ; 0 return = success
 call exit
```

```
segment .data
msg: db "Hello World!",0x0a ; String to print
len: equ $-msg ; Length of the string
```

- Think of the \$ as the current assembly address (byte addressable).
- Then \$-msg, is the current assembly address minus the address of the label msg.
- As a result len is set to equal to the number of bytes in the message.
- Better then manually counting.

# Open system call

```
int open (char *pathname, int flags [, int mode]);
```

- `pathname` is a null-terminated string
- `flags` is a collection of options or'ed together about how the file is opened.

| flags | meaning            |
|-------|--------------------|
| 0     | read-only          |
| 1     | write-only         |
| 2     | read and write     |
| 0x40  | create if needed   |
| 0x200 | truncate to size 0 |
| 0x400 | append             |

- `mode` is the permissions to grant if a file is created

# Permissions for files

- There are 3 basic permissions: read, write and execute
- There are 3 categories of users: user (owner), group and other
- Each of the 3 categories gets a 0 or 1 for each basic permission
- Octal works well for permissions
- For example 640
  - ▶ 110 100 000
  - ▶ rw- r-- ---
  - ▶ read and write permission to the user.
  - ▶ read permission to the group.
  - ▶ no permission to others.

## Code to open a file

- Open system call returns a small non-negative integer identifying the opened file (file descriptor)
- It returns -1 on error and sets errno

```
segment .data
fd: dd 0
name: db "sample",0
segment .text
extern open
lea rdi, [name] ; pathname
mov esi, 0x42 ; read-write | create
mov edx, 600o ; read-write for me
call open
cmp eax, 0
jl error ; failed to open
mov [fd], eax
```

# Read and write system calls

```
int read (int fd, void *data, long count);
int write (int fd, void *data, long count);
```

- fd is the file descriptor returned by open
- data is a pointer to some memory to send or receive data
- count is the number of bytes to read or write
- The data can be any type
- These functions return the number of bytes read or written
- read returns 0 on end-of-file
- They both return -1 on errors and set errno
- Use perror to print a text description based on errno

# Lseek system call

```
long lseek (int fd, long offset, int whence);
```

- offset is a byte offset
- If whence is 0, offset is the byte position from the start of the file
- If whence is 1, offset is relative to the current position
- If whence is 2, offset is relative to the end of the file
- lseek returns the current position
- Using whence = 2 and offset = 0, lseek returns the file size

# The close system call

```
int close (int fd);
```

- You should make a habit of closing files when no longer needed
- They will be closed when the process ends
- No data is buffered in the user process, so data written to unclosed files will not have been lost on a close call.
- Closing will reduce overhead in the kernel
- There is a per-process limit on open files
- Use “ulimit -a” to see your limits



# Example

- Let us write a simple encrypt program.
- The program will rotate each byte right by 1
- We will use the following helper functions.
  - ▶ `long file_size(int fd)`
  - ▶ `void encrypt(void* data, long size)`

## file\_size

```
;long file_size(int fd)
file_size:
 push rbp
 mov rbp, rsp
 sub rsp, 16
.fd equ 0
.size equ 4
 mov [rsp+.fd],edi
 mov rsi,0 ;offset=0
 mov rdx,2 ;whence=2
 call lseek
 mov [rsp+.size],rax
 mov edi, [rsp+.fd]
 mov rsi,0 ;offset=0
 mov rdx,0 ;whence=0
 call lseek
 mov rax,[rsp+.size]
 leave
 ret
```

# encrypt

```
;void encrypt(void* data, long size)
```

```
encrypt:
```

```
 xor rcx,rcx ; i=0
```

```
.while:
```

```
 cmp rcx,rsi ;i<size
```

```
 jge .ewhile
```

```
 mov al, [rdi+rcx] ;al=data[i]
```

```
 ror al, 1
```

```
 mov [rdi+rcx],al ;data[i]=al
```

```
 inc rcx ;++i
```

```
 jmp .while
```

```
.ewhile
```

```
 ret
```

# main program

```
 segment .data
filename: db "test.txt",0
fd: dd 0
fsize: dq 0
datap: dq 0

 segment .text
global main
extern open,lseek,malloc,read,write

main:
 push rbp
 mov rbp, rsp
```

# main program

```
mov rdi, filename
mov rsi, 2 ;read-write
call open

mov [fd], eax ; save fd
mov edi, eax
call file_size

mov [fsize],rax
mov rdi, rax
call malloc
mov [datap],rax
```

## main program

```
mov edi,[fd]
mov rsi,[datap]
mov rdx,[fsize]
call read
```

```
mov edi,[fd]
mov rsi,0
mov rdx,0
call lseek
```

```
mov rdi, [datap]
mov rsi, [fsize]
call encrypt
mov edi,[fd]
mov rsi,[datap]
mov rdx,[fsize]
call write
leave
ret
```

# Structs

# Structs

- A struct is a compound object

```
struct Customer {
 int id;
 char name[71];
 char address[71];
 int balance;
};
```

- How big is this structure?
- There is the easy, but wrong answer
- and the somewhat more complex answer



# Structs

- Let us focus on emulating a structure in assembler, without any consideration for C interfacing.
- Easy mode enabled.

# Structs

- Let us focus on emulating a structure in assembler, without any consideration for C interfacing.
- Easy mode enabled.
- An `int` is 4 bytes. We have 2 of these in our structure, so that's 8 bytes

# Structs

- Let us focus on emulating a structure in assembler, without any consideration for C interfacing.
- Easy mode enabled.
- An `int` is 4 bytes. We have 2 of these in our structure, so that's 8 bytes
- A `char` is 1 byte and we have 2 arrays of 71 characters each, so this is another 142 bytes

# Structs

- Let us focus on emulating a structure in assembler, without any consideration for C interfacing.
- Easy mode enabled.
- An `int` is 4 bytes. We have 2 of these in our structure, so that's 8 bytes
- A `char` is 1 byte and we have 2 arrays of 71 characters each, so this is another 142 bytes
- So 150 in total.

# Structs

If we want to allocate space for our struct, a simple call to `malloc` will suffice.

```
mov rdi, 150 ; size of a Customer
call malloc
mov [c], rax ; save the address
```

So `c` now holds a pointer to our allocated structure.

- But how do we use the struct?

# Structs

If we want to allocate space for our struct, a simple call to `malloc` will suffice.

```
mov rdi, 150 ; size of a Customer
call malloc
mov [c], rax ; save the address
```

So `c` now holds a pointer to our allocated structure.

- But how do we use the struct?
- By using offsets.

```
struct Customer {
 int id; [rax]----->id
 char name[71]; [rax+4]----->name
 char address[71]; [rax+75]---->address
 int balance; [rax+146]--->balance
};
```

## Filling in a C struct

```
char * strcpy (char * destination, const char * source);
```

```
segment .data
name db "Bob",0
address db "22 Duncun street",0
balance dd 123
.....
 mov [rax], dword 7 ; set the id
 lea rdi, [rax+4] ; name field
 lea rsi, [name] ; name to copy to struct
 call strcpy
 mov rax, [c]
 lea rdi, [rax+75] ; address field
 lea rsi, [address] ; address to copy
 call strcpy
 mov rax, [c]
 mov edx, [balance]
 mov [rax+146], edx
```

# Assembly struct

- Using the yasm struc pseudo-op we can define a Customer

```
 struc Customer
id resd 1
name resb 71
address resb 71
balance resd 1
 endstruc
```

- id, name, address and balance are globals
- It's almost the same as doing 4 equates
- The size is Customer\_size



# Assembly struct

- Using the yasm struc pseudo-op we can define a Customer

```
 struc Customer
id resd 1
name resb 71
address resb 71
balance resd 1
 endstruc
```

- id, name, address and balance are globals
- It's almost the same as doing 4 equates
- The size is Customer\_size
- But, you could not have id in 2 structs

# Assembly struct

- One alternative is to prefix field names with dots

```
 struc Customer
.id resd 1
.name resb 71
.address resb 71
.balance resd 1
 endstruc
```

- Then you would have to use `Customer.id`
- Another alternative is to use an abbreviated prefix

```
 struc Customer
c_id resd 1
c_name resb 71
c_address resb 71
c_balance resd 1
 endstruc
```

## Program to allocate and fill a struct - data segment

```
segment .data
name db "Calvin", 0
address db "12 Mockingbird Lane",0
balance dd 12500

struct Customer
c_id resd 1
c_name resb 71
c_address resb 71
c_balance resd 1
endstruct

c dq 0 ; to hold a Customer pointer
```

## Program to allocate and fill a struct - part of text segment

```
mov rdi, Customer_size
call malloc
mov [c], rax ; save the pointer
mov [rax+c_id], dword 7
lea rdi, [rax+c_name]
lea rsi, [name]
call strcpy
mov rax, [c] ; restore the pointer
lea rdi, [rax+c_address]
lea rsi, [address]
call strcpy
mov rax, [c] ; restore the pointer
mov edx, [balance]
mov [rax+c_balance], edx
```

# Size Discrepancy

- Now the hard question. How big would the same C++ struct be?

# Size Discrepancy

- Now the hard question. How big would the same C++ struct be?
- **152 bytes**



## C Alignment, and Struct padding

- This happens because C/C++ enforces primitives to have specific alignment based on its size.
- In effect C/C++ is padding the struct to achieve this alignment.

# C Alignment, and Struct padding

- This happens because C/C++ enforces primitives to have specific alignment based on its size.
- In effect C/C++ is padding the struct to achieve this alignment.
- Certain data types have specific alignment requirements.
- The ones relevant to us in 64-bit linux are:
  - ▶ chars(1 byte) have no alignment requirement.
  - ▶ shorts(2 bytes) must start on an even address (multiple of 2).
  - ▶ int,float(4 bytes) must start on an multiple of 4
  - ▶ long,double(8 bytes) must start on a multiple of 8
  - ▶ pointer must start on a multiple of 8
- Furthermore alignment must still be preserved across struct elements in an array.



# C Alignment, and Struct padding

- For example the struct

```
struct example
{
 char *p; // 8 bytes
 char c; // 1 byte
 int x; // 4 bytes
}
```

# C Alignment, and Struct padding

- For example the struct

```
struct example
{
 char *p; // 8 bytes
 char c; // 1 byte
 int x; // 4 bytes
}
```

- Will actually be stored as

```
struct example
{
 char *p; // 8 bytes
 char c; // 1 byte
 char pad[3]; // 3 bytes
 int x; // 4 bytes
}
```

# C Alignment, and Struct padding

- For example the struct

```
struct example
{
 char c; // 1 byte
 char *p; // 8 bytes
}
```

# C Alignment, and Struct padding

- For example the struct

```
struct example
{
 char c; // 1 byte
 char *p; // 8 bytes
}
```

- Will actually be stored as

```
struct example
{
 char c; // 1 byte
 char pad[7]; // 7 bytes
 char *p; // 8 bytes
}
```

## C Alignment, and Struct padding

- The padding also has to be applicable for aligning multiples of the same struct (arrays)

```
struct example
{
 int e; //4 bytes
 char c; //1 byte
}
```

## C Alignment, and Struct padding

- The padding also has to be applicable for aligning multiples of the same struct (arrays)

```
struct example
{
 int e; //4 bytes
 char c; //1 byte
}
```

- Will actually be stored as

```
struct example
{
 int e; //4 bytes
 char c; //1 byte
 char pad[3]; //3 bytes
}
```

# C Alignment, and Struct padding

- struct example

```
{
 long e; //8 bytes
 char c; //1 byte
}
```

## C Alignment, and Struct padding

- struct example

```
{
 long e; //8 bytes
 char c; //1 byte
}
```

- Will actually be stored as

```
struct example
{
 long e; //8 bytes
 char c; //1 byte
 char pad[7]; //7 bytes
}
```



# C Alignment, and Struct padding

- struct example

```
{
 int a; //4 bytes
 long b; //8 bytes
 char c; //1 byte
}
```

# C Alignment, and Struct padding

- struct example

```
{
 int a; //4 bytes
 long b; //8 bytes
 char c; //1 byte
}
```

- Will actually be stored as

```
struct example
{
 int a; //4 bytes
 char pad[4]; //4 bytes
 long b; //8 bytes
 char c; //1 byte
 char pad2[7]; //7 bytes
}
```

Why 7 and not 3?

## Allocating a slightly more complex array of customers

```
 segment .data
 struc Customer
c_id resd 1 ; 4 bytes in total
c_name resb 65 ; 69 bytes in total
c_address resb 65 ; 134 bytes in total
 align 4 ; aligns to 136
c_balance resd 1 ; 140 bytes in total
c_rank resb 1 ; 141 bytes in total
 align 4 ; aligns to 144
 endstruc
customers dq 0
 segment .text
 mov rdi, 100 ; for 100 structs
 imul rdi, Customer_size
 call malloc
 mov [customers], rax
```

# Printing an array of customers

```
segment .data
format db "%s %s %d",0x0a,0
segment .text
push rbp
mov rbp, rsp
push r15
push r14
mov r15, 100 ; counter saved through calls
mov r14, [customers]; pointer saved through calls
more lea rdi, [format]
 lea rsi, [r14+c_name]
 lea rdx, [r14+c_address]
 mov ecx, [r14+c_balance]
 mov rax, 0
 call printf
 add r14, Customer_size
 dec r15
 jnz more
 pop r14
 pop r15
 leave
```

# Data Structures

# Data structures


- In this lecture we will not learn any \*new\* assembler instructions/concepts.
- But rather, we will utilize the knowledge we have gained so far to build some common data structures.
  - ▶ One obvious downside to using ASM to implement a data structure, is that generality via templatization is not possible. (unless you write part of a compiler for your own tweaked version of assembler)
  - ▶ However given ASM's low level nature we can utilize the specific details of the intended type to our advantage. Namely in terms of
    - ★ Speed
    - ★ Size

# Outline

The data structures we will be focusing on are as follows

- Linked Lists
- Double Link Lists
- Hash Tables
- Binary and K-Array trees
- Directed Graphs

# Linked lists



- A simple linked list is constructed of a sequence of structs
- Each struct has some data and a pointer to the next item on the list
- The filled circle means a pointer equal to NULL (0)
- There needs to be some memory cell containing the first pointer
- This list has no obvious order to the keys
- It could be ordered by insertion time in two ways: by inserting at the front or the end
- It is easier to insert at the front, though the value of `list` will change with each insertion



# List node struct definition

```
 struc node
n_value resq 1
n_next resq 1
 endstruc
```

# Creating an empty list

- The only requirement will be to set the pointer to NULL
- Having a function makes it possible to change later with less impact on the rest of the program

newlist:

```
 mov rax, 0
```

```
 ret
```

```
 ...
```

```
 call newlist
```

```
 mov [list], rax ; where list is a dq
```

# Inserting a number into a list

- A new node will be allocated and placed at the start
- We must pass the list pointer into the function
- We also must receive a new pointer back to store in `list`
- In C++ the prototype would be `node* insert(node*, long)` and we would use

```
list = insert (list, k);
```

- In assembly we would insert `k` using

```
mov rdi, [list] ; pass in the list pointer
mov rsi, [k]
call insert
mov [list], rax ; we have a new list pointer
```

## Insert code

```
insert:
.list equ 0
.k equ 8
 push rbp
 mov rbp, rsp
 sub rsp, 16
 mov [rsp+.list], rdi ; save list pointer
 mov [rsp+.k], rsi ; and k on stack
 mov rdi, node_size
 call malloc ; rax will be node pointer
 mov r8, [rsp+.list] ; get list pointer
 mov r9, [rsp+.k] ; get k
 mov [rax+n_next], r8 ; new node point to old list
 mov [rax+n_value], r9; save k in node
 leave
 ret
```

# Printing the list

```
;print(node* list)
print:
 push rbp
 mov rbp, rsp
 push r12
 mov r12,rdi ; save list
 ; we use r12 since since it
 ; is safe across function calls

 segment .data
.prt_frm: db "%ld ",0
.prt_nl: db 0x0a,0
```

## Printing the list cont.

```
segment .text
 .while:
 cmp r12,0
 je .ewhile
 mov rdi, print.prt_frm
 mov rsi, [r12+n_value]
 xor eax,eax
 call printf
 mov r12,[r12+n_next]
 jmp .while
 .ewhile:

 mov rdi, print.prt_nl
 xor eax,eax
 call printf

 pop r12
 leave
 ret
```

# Main program to build a list, part 1

```
main:
.list equ 0
.k equ 8
segment .data
 scanf_fmt: db "%ld",0
segment .text
```

## Main program to build a list, part 2

```
 push rbp
 mov rbp, rsp
 sub rsp, 16
 call newlist
 mov [rsp+.list], rax ; .list equal to 0
.more lea rdi, [.scanf_fmt] ; .scanf_fmt - "%ld",0
 lea rsi, [rsp+.k] ; .k equal to 8
 xor eax, eax ; no floating point value parameters
 call scanf
 cmp rax, 1 ; quit if scanf does not return 1
 jne .done
 mov rdi, [rsp+.list] ; Get the list pointer
 mov rsi, [rsp+.k] ; Get k
 call insert
 mov [rsp+.list], rax ; Save new list pointer
 mov rdi, rax ; Move the pointer to be a parameter
 call print
 jmp .more ; Try to read another number
.done leave
 ret
```



# Example


```
4
4
5
5 4
7
7 5 4
8
8 7 5 4
1
1 8 7 5 4
2
2 1 8 7 5 4
6
6 2 1 8 7 5 4
4
4 6 2 1 8 7 5 4
1
1 4 6 2 1 8 7 5 4
2
2 1 4 6 2 1 8 7 5 4
f
```

# Ordered Insert

## Homework:

- Implement a deallocate method for the list.
- Implement an insert at the end of the list.
- Implement an ordered insert, using insert as a helper function (where we assume the list is ordered before the call).

# Doubly linked lists




- This list uses forwards and backwards pointers to make a cycle
- Also the first node is not used, so an empty list will have one node and will be circular
- The first node is called a “head” node
- Using a head node and a circular list makes insertion trivial
- You can also insert and remove from either end easily

# Doubly linked list node struct

```
 struc node
n_value resq 1
n_next resq 1
n_prev resq 1
 endstruc
```


- An “empty” list is still circular
- There are no special cases to consider



## Doubly linked list node :new list

```
; list=newlist()
newlist:
 push rbp
 mov rbp, rsp
 mov rdi, node_size
 call malloc
 mov [rax+n_next], rax
 mov [rax+n_prev], rax
 leave
 ret
```

# Inserting at the front of a doubly linked list



- Make the new node's `n_next` point to the head cell's next
- Make the new node's `n_prev` point to the head cell
- Make the head's next cell point to the new cell
- Make the new node's next node's `prev` cell point backward to the new node

The original links are dashed lines

# Insertion function

```
; insert (node* list, long k);
insert: push rbp
 mov rbp, rsp
 sub rsp, 16
 mov [rsp+.list], rdi ; save list pointer, .list equ 0
 mov [rsp+.k], rsi ; and k on stack, .k equ 8
 mov rdi, node_size
 call malloc ; rax will be node pointer
 mov r8, [rsp+.list] ; get list pointer
 mov r9, [r8+n_next] ; get head's next
 mov [rax+n_next], r9 ; set new node's next
 mov [rax+n_prev], r8 ; set new node's prev
 mov [r8+n_next], rax ; set head's next
 mov [r9+n_prev], rax ; set new node's next's prev
 mov r9, [rsp+.k] ; get k
 mov [rax+n_value], r9 ; save k in node
 leave
 ret
```


# List traversal

```
; print (node* list);
print: push rbp
 mov rbp, rsp
 sub rsp, 16
 mov [rsp+.r12], r12 ; save r12, .r12 equ 0
 mov [rsp+.list], rdi ; save list, .list equ 8
 mov r12, [rdi+n_next] ; skip the head node
 cmp r12, [rsp+.list] ; is the list empty? (head.next==head)
 je .done
.more lea rdi, [.print_fmt] ; .print_fmt: db "%ld ",0
 mov rsi, [r12+n_value]
 xor eax, eax
 call printf ; print the node's value
 mov r12, [r12+n_next] ; advance to the next node
 cmp r12, [rsp+.list] ; have we reached the head cell?
 jne .more
.done lea rdi, [.newline] ; .newline: db 0x0a,0
 xor eax, eax
 call printf
 mov r12, [rsp+.r12] ; restore r12
 leave
 ret
```



# Hash tables

- For each key, compute a hash value.  $A[hash(key)]$
- The hash value defines an index in an array to store a value associated with a given key.
- Collisions occur when 2 different keys hash to the same index
- The simplest collision resolution is to use a linked list



# A hash function for integers

- A good hash function spreads the keys around
- Using  $(k \bmod t)$  as a hash, where  $t$  is the table size, is sufficient for our purposes
- A good recommendation is to make  $t$  prime, but
- in this example,  $t = 256$ , so using **and** works

```
; i = hash (n);
hash mov rax, rdi
 and rax, 0xff
 ret
```

- For example  $-4 \rightarrow 111..100 \ \& \ 0xff = 0..011111100 \rightarrow 252$

# A hash function for strings

- The code below uses the characters of the string as coefficients of a polynomial
- The polynomial is evaluated at 191 (a prime)
- Then a mod is done with 100000 to get the hash value
- Homework: write this in assembler. We use Horner's Method to calculate the polynomial quicker.

```
int hash (unsigned char *s)
{
 unsigned long h = 0;
 int i = 0;
 while (s[i]) {
 h = h*191 + s[i];
 i++;
 }
 return h % 100000;
}
```

# Hash node structure and array of pointers

- Our hash table that uses integer keys has only 256 entries/pointers
- Usually the array would be larger and a creation function needed

```
segment .data
table times 256 dq 0 ; All NULL pointers
 struc node
n_value resq 1
n_next resq 1 ; Singly linked list
 endstruc
```

## Function to find a key

```
; p = find (n); where p is a pointer to the matched node.
; p = 0 if not found
find: push rbp
 mov rbp, rsp
 sub rsp, 16
 mov [rsp], rdi ; save key
 call hash
 mov rax, [table+rax*8] ; get pointer
 mov rdi, [rsp] ; get key
 cmp rax, 0 ; empty list?
 je .done
.more cmp rdi, [rax+n_value] ; key match?
 je .done
 mov rax, [rax+n_next] ; advance on the collision list
 cmp rax, 0 ; end of list
 jne .more
.done leave
 ret
```

## Function to insert a key

```
insert: push rbp
 mov rbp, rsp
 sub rsp, 16
 mov [rsp+.n], rdi ; save n, .n equ 0
 call find
 cmp rax, 0 ; Is n already there?
 jne .found
 mov rdi, [rsp+.n] ; compute hash(n)
 call hash
 mov [rsp+.h], rax ; save hash value, .h equ 8
 mov rdi, node_size ; allocate a node
 call malloc
 mov r9, [rsp+.h] ; use r9 as index register
 mov r8, [table+r9*8] ; get old pointer from table
 mov [rax+n_next], r8 ; make new node point to old
 mov r8, [rsp+.n] ; get n from the stack
 mov [rax+n_value], r8 ; set the node value
 mov [table+r9*8], rax ; make new node first on its list
.found leave
 ret
```

# Binary trees

- A binary tree is a hierarchy of nodes
- There is a root node (or not, for an empty tree)
- Each node can have a left child and a right child
- The node structure has 2 pointers
- Either or both pointers could be NULL
- Binary trees usually impose an ordering on the nodes.
  - ▶ We will assume keys less than the current node's key will be stored down the left branch
  - ▶  $Leftchild < Current < Rightchild$
- Such a tree is a “binary search tree”
- We do not store duplicate keys.

# Node Structure

```
 struc node
n_value resq 1
n_left resq 1
n_right resq 1
 endstruc
```



# A structure for the tree

- We could represent an empty tree as a NULL pointer
- This introduces special cases
  - ▶ specifically any operation that may change the position of the root node, would need to return an updated pointer to the tree(new root node).
- Instead we implement a tree struct
- It contains the root pointer which can be NULL
- It also contains the count of nodes in the tree
- After creating a tree, we use the same pointer for all function calls

```
 struc tree
t_count resq 1
t_root resq 1
 endstruc
```

## Creating a new tree

- The `new_tree` function allocates a tree struct and sets it up as an empty tree

`new_tree:`

```
 push rbp
 mov rbp, rsp

 mov rdi, tree_size
 call malloc
 xor edi, edi
 mov [rax+t_root], rdi
 mov [rax+t_count], rdi

 leave
 ret
```

## Finding a node in a tree: `p = find(tree,n)`

```
find: push rbp
 mov rbp, rsp
 mov rdi, [rdi+t_root]
 xor eax, eax
.more cmp rdi, 0
 je .done
 cmp rsi, [rdi+n_value]
 jl .goleft
 jg .goright
 mov rax, rdi ; note, in the book this is rsi
 jmp .done
.goleft:
 mov rdi, [rdi+n_left]
 jmp .more
.goright:
 mov rdi, [rdi+n_right]
 jmp .more
.done leave
ret
```

# Inserting a node into a tree

## Overview:

- First you check to see if the key is already in the tree
- If not, then you create a new node and set its value to the given key and set its two children to NULL.
- There is a special case for an empty tree
  - ▶ specifically, you make the newly created node the root of the tree
- If the tree is not empty, then we must traverse down the tree, going sometimes left and sometimes right to find the correct place to insert the new node

## Inserting a node into a tree: `p=insert(tree,n)`

`insert:`

```
.n equ 0
.t equ 8

 push rbp
 mov rbp, rsp
 sub rsp, 16
 mov [rsp+.t], rdi
 mov [rsp+.n], rsi
 call find
 cmp rax, 0
 jne .done
 mov rdi, node_size
 call malloc
 mov rsi, [rsp+.n]
 mov [rax+n_value], rsi
 xor edi, edi
 mov [rax+n_left], rdi
 mov [rax+n_right], rdi
```

# Inserting a node into a tree

```
mov rdx, [rsp+.t]
mov rdi, [rdx+t_count]
cmp rdi, 0
jne .findparent
inc qword [rdx+t_count]
mov [rdx+t_root], rax
jmp .done
```

# Inserting a node into a tree

.findparent:

```
 inc qword [rdx+t_count]
 mov rdx, [rdx+t_root]
```

.repeatfind:

```
 cmp rsi, [rdx+n_value] ; Remember rsi=n
 jl .goleft
 mov r8, rdx
 mov rdx, [r8+n_right]
 cmp rdx, 0
 jne .repeatfind
 mov [r8+n_right], rax
 jmp .done
```

.goleft:

```
 mov r8, rdx
 mov rdx, [r8+n_left]
 cmp rdx, 0
 jne .repeatfind
 mov [r8+n_left], rax
```

.done leave

ret

# Printing the keys in order

- We first call a non-recursive function with the tree object
- It calls a recursive function with the root node

```
; print(t);
print:
 push rbp
 mov rbp, rsp
 mov rdi, [rdi+t_root]
 call rec_print

segment .data
.print db 0x0a, 0

segment .text
 lea rdi, [.print]
 xor eax, eax
 call printf
 leave
 ret
```



## Recursive print function: rec\_print(t)

```
rec_print: push rbp
 mov rbp, rsp
 sub rsp, 16 ; make room to save t
 cmp rdi, 0 ; return if t is NULL
 je .done
 mov [rsp+.t], rdi ; save t, .t equ 0
 mov rdi, [rdi+n_left] ; print the left sub-tree
 call rec_print
 mov rdi, [rsp+.t] ; print the current node
 mov rsi, [rdi+n_value]
 lea rdi, [.print2] ; .print2: bd "%ld ",0
 xor rax, rax
 call printf
 mov rdi, [rsp+.t] ; print the right sub-tree
 mov rdi, [rdi+n_right]
 call rec_print
.done leave
 ret
```

# Homework

- Deallocate a binary tree
- Print the content of the tree layer by layer (Breadth first traversal).  
*Remember you actually have a stack available*

## High Performance ASM

# High Performance ASM

We will now discuss a couple of common and effective ways of optimizing ASM (or code in general).

- But first, it should be noted, that it is not always viable to write the most efficient code for every part of a system. Rather spend the time and effort on the bottlenecks in the system.

# High Performance ASM

We will now discuss a couple of common and effective ways of optimizing ASM (or code in general).

- But first, it should be noted, that it is not always viable to write the most efficient code for every part of a system. Rather spend the time and effort on the bottlenecks in the system.



# Beating the compiler

- An optimizing compiler will implement nearly all of the general optimizations
- It will do them tirelessly, missing nearly nothing
- You must usually find a non-obvious technique to get better performance than the compiler
- Learn from the compiler
  - ▶ Use the `-S` option to get an assembly listing
  - ▶ Perhaps you can do the compiler's tricks better


# Efficient use of cache

- Our aim is generally to keep a processor as busy as possible.
- If a CPU is working with cached data it doesn't wait around much.
- If the required data is in memory and not in the cache, a substantial amount of computational capacity is wasted while the data is retrieved.
- Organize your algorithm to exploit data locality.

# Efficient use of cache: Example

- The plot below shows time versus array size for computing 10 billion exclusive or operations on quad words.

Time to Compute XOR





# Efficient use of cache: Example

- Consider the task of matrix multiplication (with square matrices for simplicity)



$$\begin{bmatrix} a_{11} & a_{12} & \dots & a_{1N} \\ a_{21} & a_{22} & \dots & a_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ a_{N1} & a_{N2} & \dots & a_{NN} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} & \dots & b_{1N} \\ b_{21} & b_{22} & \dots & b_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ b_{N1} & b_{N2} & \dots & b_{NN} \end{bmatrix} \quad (1)$$

$$= \begin{bmatrix} \sum_{k=1}^N a_{1k} b_{k1} & \sum_{k=1}^N a_{1k} b_{k2} & \dots & \sum_{k=1}^N a_{1k} b_{kN} \\ \sum_{k=1}^N a_{2k} b_{k1} & \sum_{k=1}^N a_{2k} b_{k2} & \dots & \sum_{k=1}^N a_{2k} b_{kN} \\ \vdots & \vdots & \ddots & \vdots \\ \sum_{k=1}^N a_{Nk} b_{k1} & \sum_{k=1}^N a_{Nk} b_{k2} & \dots & \sum_{k=1}^N a_{Nk} b_{kN} \end{bmatrix} \quad (2)$$

# Efficient use of cache: Example

- Naive matrix multiplication can easily be implemented with three nested loops, and is  $O(N^3)$

```
for(int i=0; i<N; ++i)
 for(int j=0; j<N; ++j)
 {
 sum = 0.0;
 for(int k=0; k<N; ++k)
 sum+= a[i][k] * b[k][j]
 c[i][j]= sum;
 }
```

# Efficient use of cache: Example


- Naive matrix multiplication can easily be implemented with three nested loops, and is  $O(N^3)$

```
for(int i=0; i<N; ++i)
 for(int j=0; j<N; ++j)
 {
 sum = 0.0;
 for(int k=0; k<N; ++k)
 sum+= a[i][k] * b[k][j]
 c[i][j]= sum;
 }
```

- There are three indices that are being iterated over. Which means there are at most  $\binom{3}{2} = 6$  possible unique ways of iterating.
- Which way is best? We will consider three cases (which covers 6 performance profiles)


# Matrix multiplication: IJK

```
for(int i=0; i<N; ++i)
 for(int j=0; j<N; ++j)
 {
 sum = 0.0;
 for(int k=0; k<n; ++k)
 sum+= a[i][k] * b[k][j];
 c[i][j]= sum;
 }
```




$(i, *)$

A



$(*, j)$

B




$(i, j)$


C

# Matrix multiplication: JKI


```
for(int j=0; j<N; ++j)
 for(int k=0; k<n; ++k)
 {
 r = b[k][j];
 for(int i=0; i<N; ++i)
 c[i][j] += a[i][k]*r;
 }
```



$(*,k)$   
A




$(k,j)$   
B




$(*,j)$   
C

# Matrix multiplication: IKJ


```
for(int i=0; i<N; ++i)
 for(int k=0; k<n; ++k)
 {
 r =a[i][k];
 for(int j=0; j<N; ++j)
 c[i][j] += r*b[k][j];
 }
```



(i,k)  
A



(k,\*)  
B



(i,\*)  
C

# Common sub-expression elimination

- In compiler theory, common subexpression elimination (CSE) is a compiler optimization that searches for instances of identical expressions (i.e., they all evaluate to the same value), and analyses whether it is worthwhile replacing them with a single variable holding the computed value.
- Simple example

```
a = b * c + g;
```

```
d = b * c * e;
```

could be replaced with

```
tmp = b * c;
```

```
a = tmp + g;
```

```
d = tmp * e;
```

# Strength reduction

- The idea is to rely on simpler and less computationally expensive operations.
- For example
  - ▶ Instead of dividing an integer by 8 we could shift it right 3 bits.
  - ▶ Rather than using `pow(x,3)` use `x * x * x`
  - ▶ Compute  $x^4$  by computing  $x^2$  and then squaring that
  - ▶ Avoid division by a floating point number  $x$ , rather computing  $1/x$  and use multiplication instead
  - ▶ Getting a remainder after division by 1024, can be done using **and**.

$x \text{ and } (1024 - 1)$

- It should be noted that sometimes the computationally cost of an operation is architecture dependent.



# Use registers efficiently

- Place commonly-used values in registers
- When a value is used frequently it will be better to store it in a register as apposed to in memory.

# Use fewer branches

- The problem with branches is they interrupt the instruction pipeline. Specifically it means a predictive pipelining might have to be undone.
- The compiler will frequently re-order blocks of code to reduce branches
- This is where conditional moves can help allot.

## Convert loops to branch at the bottom

- The compiler generally does this to reduce the number of instructions in a loop and, especially, the number of branches
- Here is a C for loop

```
for (i = 0; i < n; i++)
{
 x[i] = a[i] + b[i];
}
```

- By adding an **if** at the start you can loop with a branch at the bottom
- Don't do this in C or C++ (the compiler can handle this)

```
if (n > 0)
{
 i = 0;
 do {
 x[i] = a[i] + b[i];
 i++;
 } while (i < n);
}
```

# Unroll loops

- Use `-funroll-all-loops` to have gcc unroll loops
- Unrolling means repeated occurrences of the loop body with multiple parts of the data being processed
- Try to make each unrolling use different registers to reduce instruction dependence
- This frees up the CPU to do out-of-order execution
- It can do more pipelining and more parallel execution

# Unroll loops

- Consider the following function that sums the element of quad words

```
add_array: ; long add_array(long* array, long size)
```

```
 xor rax,rax
```

```
 .add_qwords
```

```
 add rax,[rdi]
```

```
 add rdi, 8
```

```
 dec rsi
```

```
 jg .add_qwords
```

```
 ret
```

# Unroll loops

- For simplicity let's consider a multiple of 4 sized array

```
add_array: long add_array(long* array, long size)
```

```
 xor rax,rax
```

```
 xor r8,r8
```

```
 xor rcx,rcx
```

```
 xor rdx,rdx
```

```
 .add_qwords:
```

```
 add rax, [rdi]
```

```
 add r8, [rdi+8]
```

```
 add rcx, [rdi+16]
```

```
 add rdx, [rdi+24]
```

```
 add rdi, 32
```

```
 sub rsi, 4
```

```
 jg .add_words
```

```
 add rcx, rdx
```

```
 add rax, r8
```

```
 add rax, rcx
```

```
 ret
```

# Merge loops

- If 2 loops have some loop limits, consider merging the bodies
- There will be less loop overhead
- The following 2 loops can be profitably merged

```
for (i = 0; i < 1000; i++) a[i] = b[i] + c[i];
for (j = 0; j < 1000; j++) d[j] = b[j] - c[j];
```

- After merging values for b[i] and c[i] can be used twice

```
for (i = 0; i < 1000; i++)
{
 a[i] = b[i] + c[i];
 d[i] = b[i] - c[i];
}
```

# Split loops

- If a loop is running over two disjoint pieces of data it might be profitable to split the loops.
- The potential cache usage can add a greater speed increase than the extra overhead incurred by having two loops.
- Real world testing is often the best way to determine if the split is worth it.



# Move loop-invariant code outside the loop

- Simple example

```
for(i=0;i<N;++i)
{
 myarray[i]*=pow (4.5,2.5)
}
```

Now if this is executed exactly as stated, the pow function will be called  $N$  times. Rather just compute it once, as reuse the result.

# Remove recursion

- Eliminating tail-recursion is generally useful
- If you have to simulate a “stack” like recursion gives you, recursion will probably be faster

# Eliminate stack frames

- Use `-fomit-frame-pointers` with `gcc`
- Use this for release code
- Using the `rbp` register is optional
- Leaf functions don't even need to worry about stack alignment
  - ▶ Unless you are using some local data requiring 16 byte alignment

# Inline functions

- The compiler can do this painlessly
- In assembly you will make your code less readable
- Explore using macros

# Reduce dependencies to allow super-scalar execution

- Use different registers to try to reduce dependencies
- The CPU has multiple computational units in 1 core
- You can benefit from out-of-order execution
- You can get more out of pipelines
- You can keep more computational units busy
- For example from our earlier `add_array` function

```
add rax, [rdi]
add r8, [rdi+8]
add rcx, [rdi+16]
add rdx, [rdi+24]
```

The order of execution does not matter, so super-scalar execution is possible.

# Use specialized instructions

- The compiler will have a harder time doing this than you
- There are also SIMD integer instructions
- Use SIMD floating point instructions
- The AVX instructions are a new feature which allow twice as many floating point values in the SIMD registers

# Use specialized instructions

Remember the question: Assume A and B have been defined in the data segment as quad words. Write the necessary ASM to calculate the hamming distance between A and B and store the result in the rax register.

- Using special instruction this is very easy

# Use specialized instructions

Remember the question: Assume A and B have been defined in the data segment as quad words. Write the necessary ASM to calculate the hamming distance between A and B and store the result in the rax register.

- Using special instruction this is very easy
- ```
mov rax, [A]  
xor rax, [B]  
popcnt rax,rax
```